

# Neural network applications



Plan for today:

- Introduce 3 deep learning research projects
- Investigate issues of bias, ethics, and fairness in AI.

# **ALGORITHMS FOR ROBUST LIFELONG LEARNING**



**Shixian Wen, Amanda Rios, Yunhao Ge,  
Laurent Itti**

This work was supported in part by C-BRIC, one of six centers in JUMP, a Semiconductor Research Corporation (SRC) program sponsored by DARPA

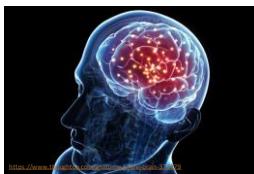
# Lifelong Learning example: human brain

Day 1



- Human brain learns many tasks from one or more domains over its lifetime.

Day 2



- Each day the human brain is exposed to new tasks, it should learn about them

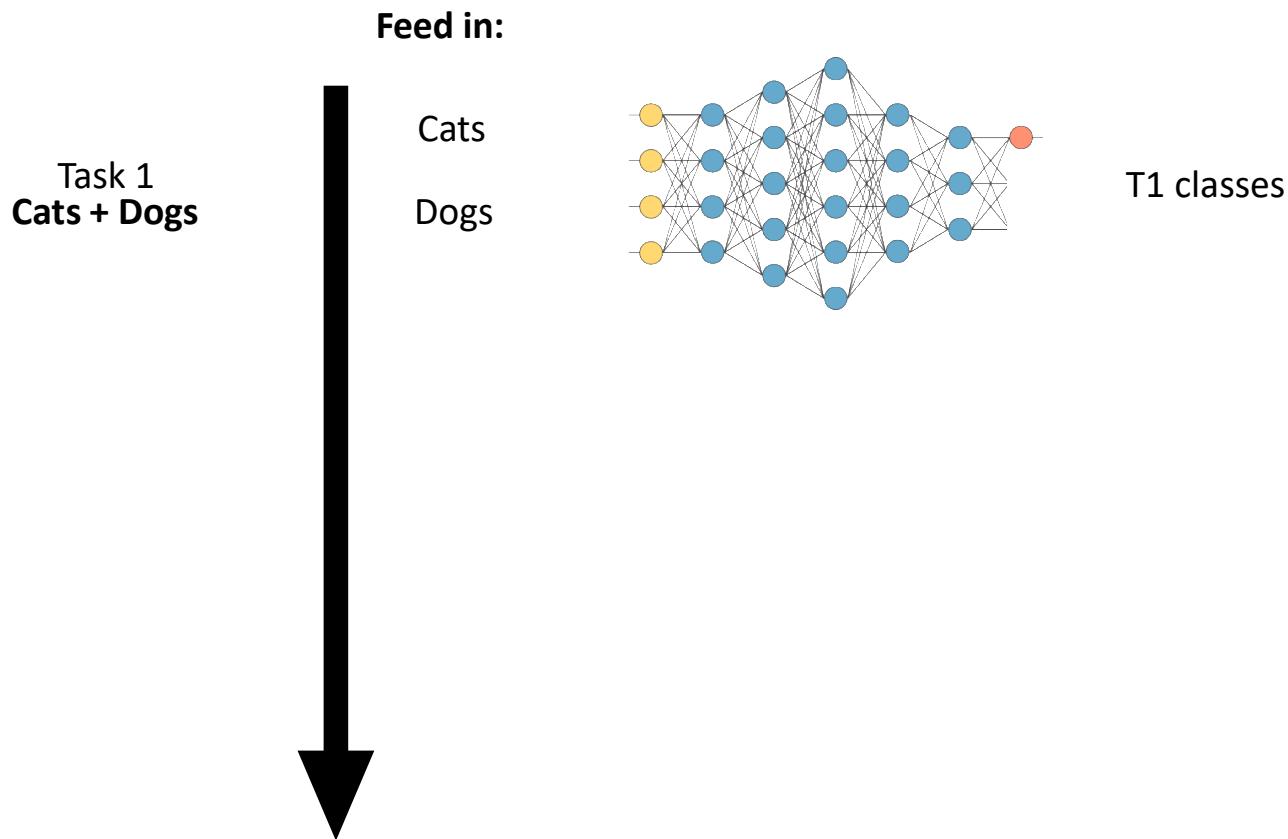
...

Day N

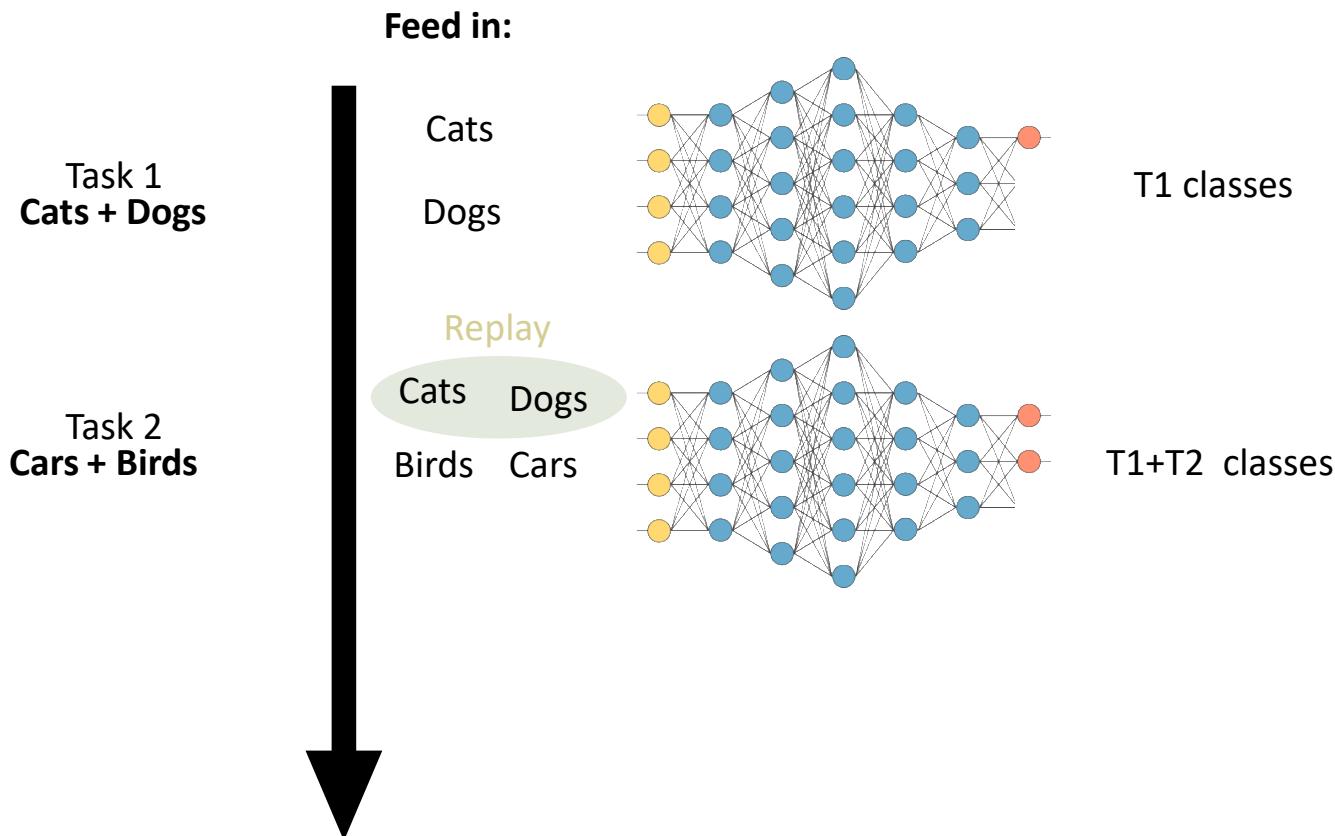


- The difficult part:
  - Should not forget tasks it has previously learned about
  - Selectively transfer the learned knowledge when learning a new task
  - Build a system that accumulatively learn about new things that we encounter everyday

# Incremental learning



# Incremental learning



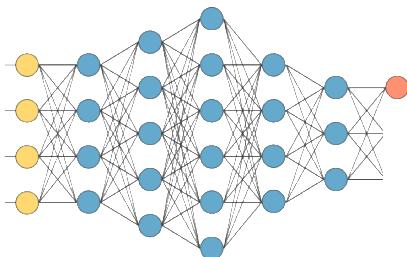
# Incremental learning

Task 1  
**Cats + Dogs**

Feed in:

Cats

Dogs

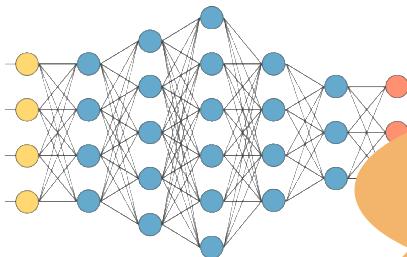


Very memory & compute expensive

Task 2  
**Cars + Birds**

Replay

Cats Dogs  
Birds Cars



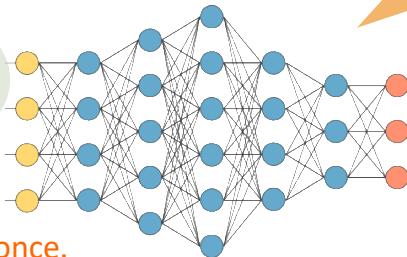
Cannot learn tasks sequentially and benefit from the training of old tasks

Task 3  
**Books + Trees**

Replay

Cats Dogs  
Cars Birds

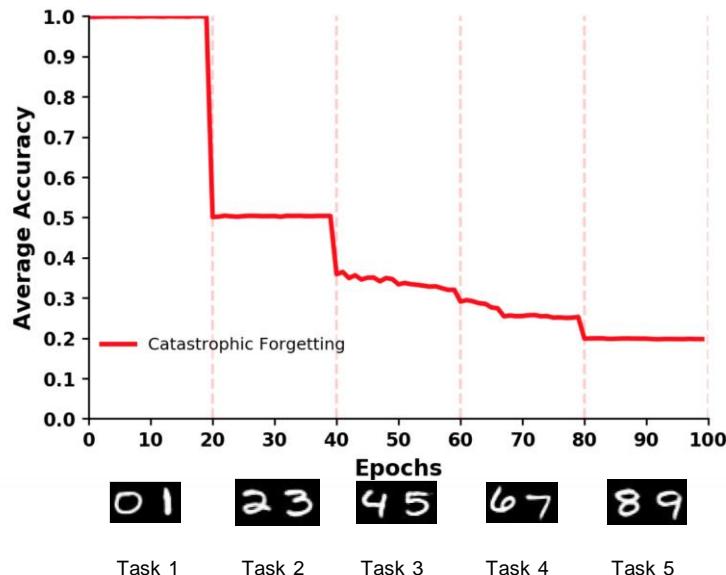
Books Trees



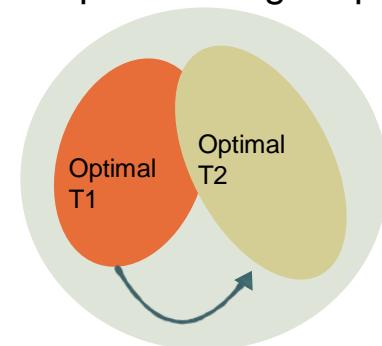
T1 + T2 + T3 classes

e.g., ImageNet: train on all 1,000 categories at once.

## Catastrophic forgetting (McCloskey & Cohen, 1989)



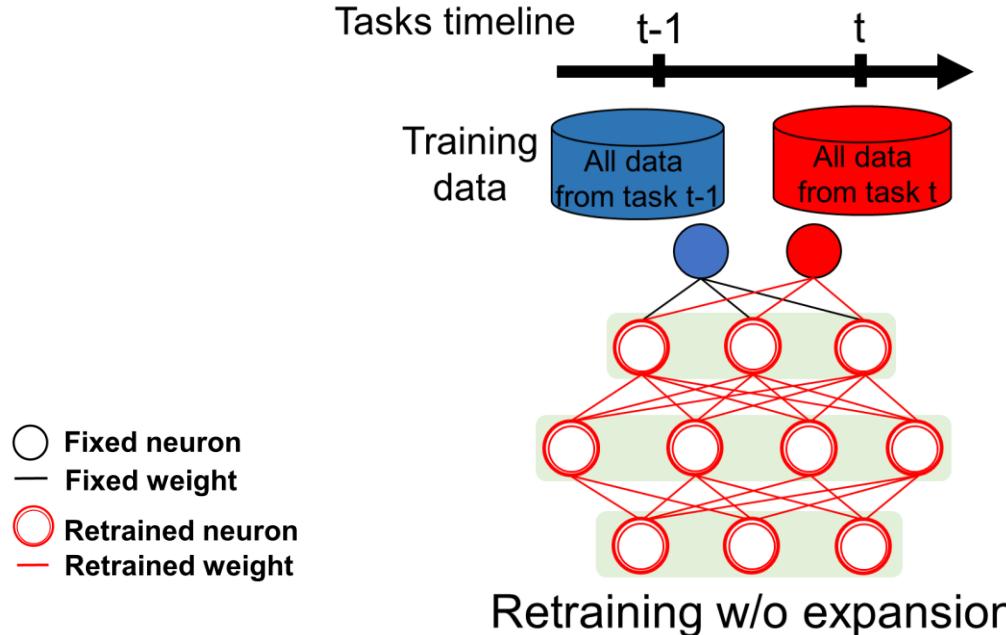
Optimal Weight Space:



Loss of supervision from T1 no longer guarantees  $W$  will be in the optimal space of T1

## Four classes of approaches: type 1

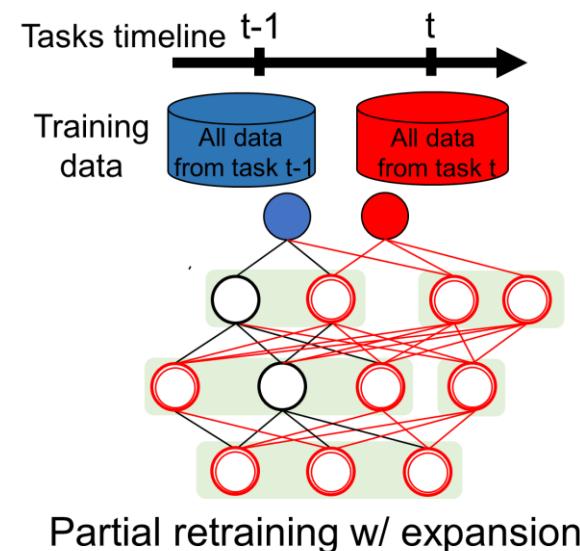
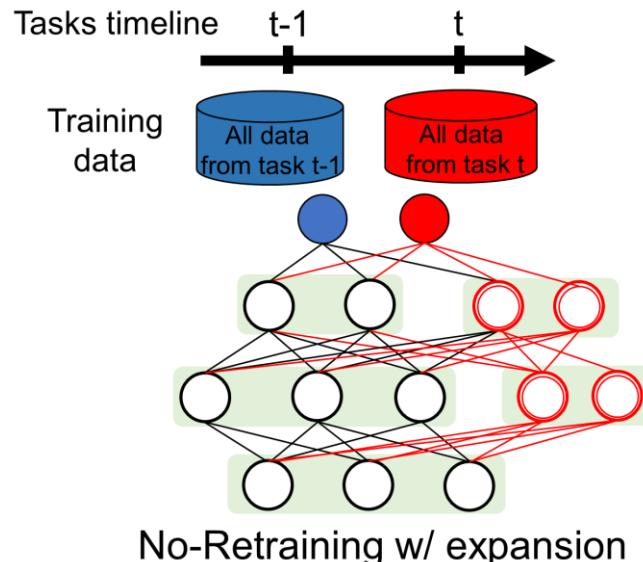
### Type 1: constrain the network weights for previous tasks



e.g., elastic weight consolidation (EWC),  
Kirkpatrick et al., 2017

## Four classes of approaches: type 2

**Type 2: expand the network capacity for new task and constrain the network weights for previous tasks**

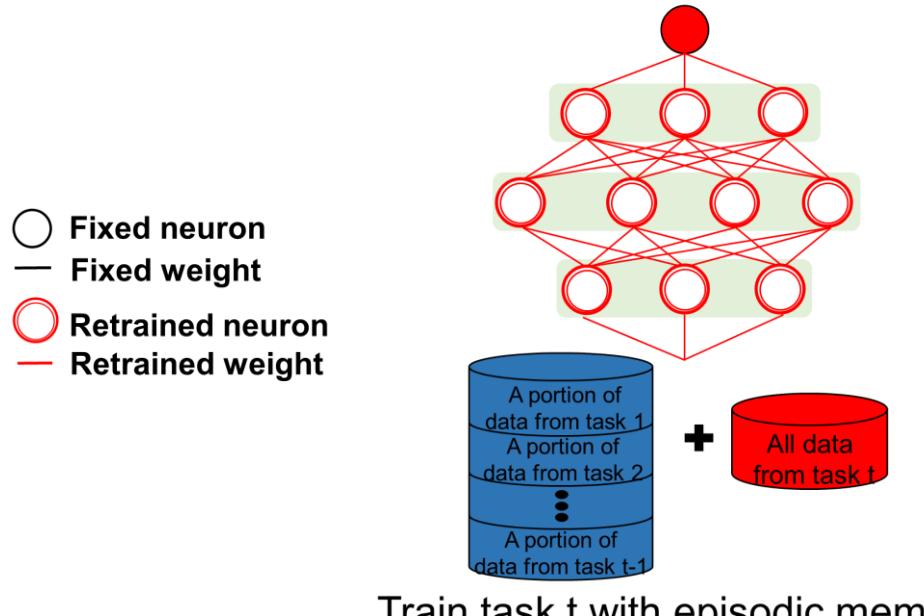


- Fixed neuron
- Fixed weight
- Retrained neuron
- Retrained weight

e.g., progressive neural nets,  
dynamically expandable nets,  
Rusu et al., 2017; Yoon et al., 2017

## Four classes of approaches: type 3

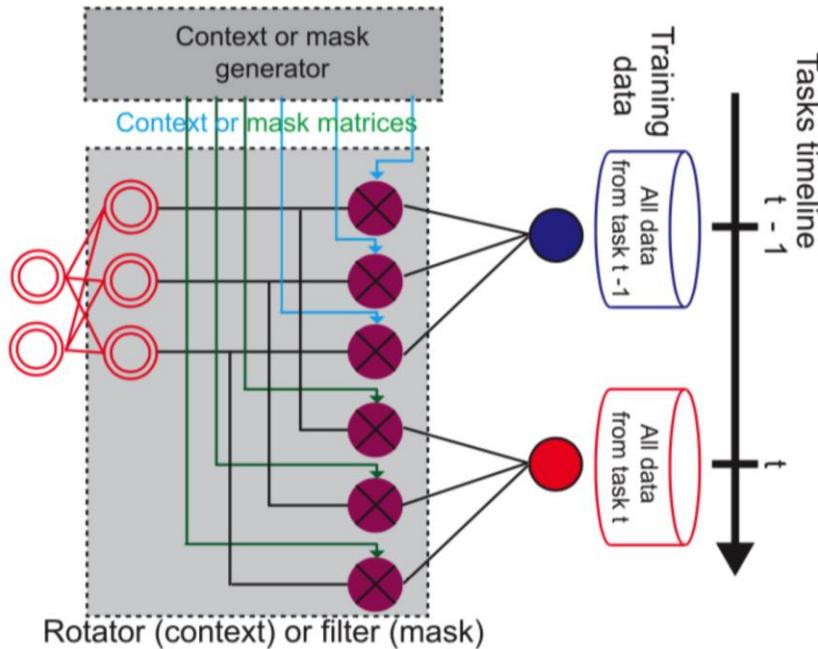
### Type 3: store a subset of original dataset



e.g., Incremental classifier and representation learning (ICARL),  
Gradient episodic memory (GEM),  
Rebuffi et al., 2017, Lopez-Paz et al., 2017

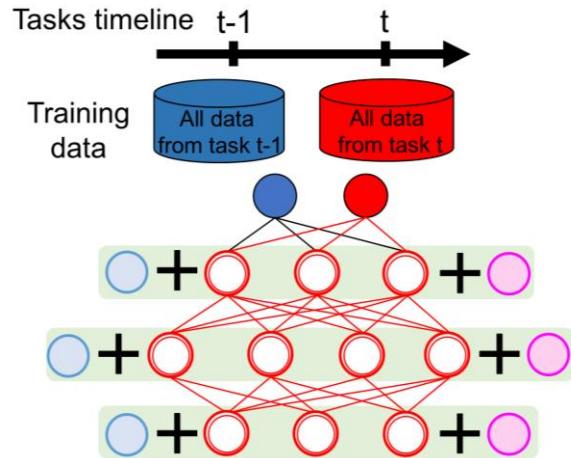
## Four classes of approaches: type 4

Type 4: Using context or mask matrices to create several subspaces or sub-networks in a big network for different tasks



e.g., parameter super position (PSP), Cheung et al., 2019

## Four classes of approaches: type 5

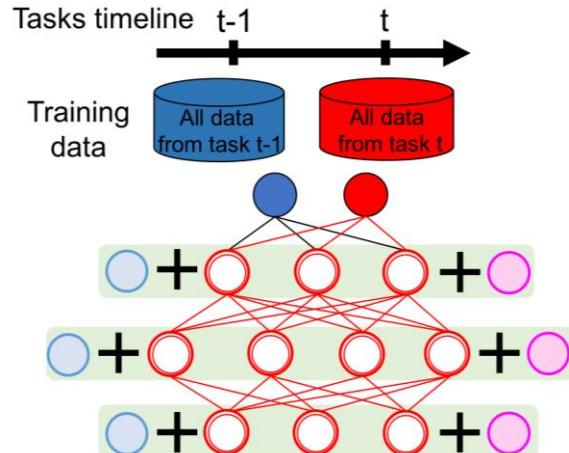


Partial retraining w/ beneficial perturbations  
Stored in the bias units

- Fixed neuron
- Fixed weight
- Retrained neuron
- Retrained weight
- Bias units for task t-1
- Bias units for task t

New proposed approach,  
Wen et al., 2021

## Four classes of approaches: type 5



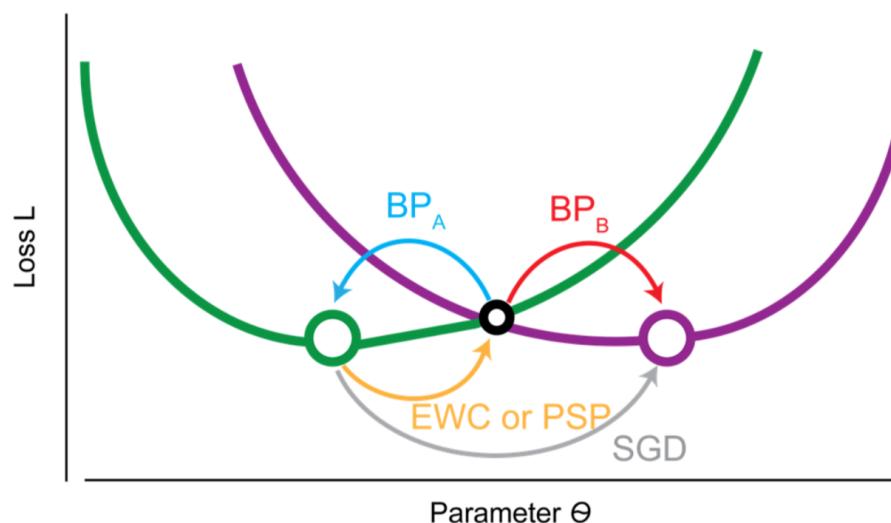
- Fixed neuron
- Fixed weight
- Retrained neuron
- Retrained weight
- Bias units for task t-1
- Bias units for task t

	Type 1	Type 2	Type 3	Type 4	Type 5 (ours)
Parameter efficient	✓	✗	✓	✗ / ✓	✓
Memory efficient	✓	✓	✗	✓	✓
Does not require previous task data	✓	✓	✗	✓	✓
Multiple input to output mappings	✗	✗	✗	✓	✓
unlimited capacity to accommodate new tasks	✓	✓	✓	✗	✓

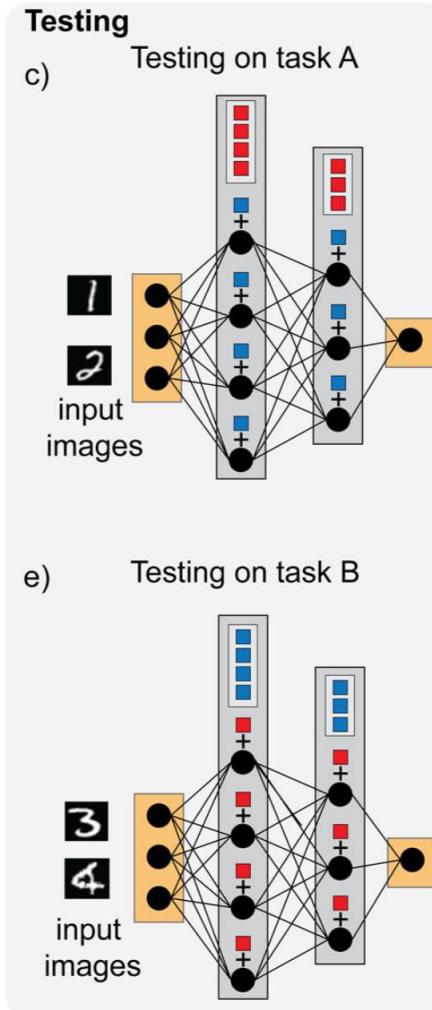
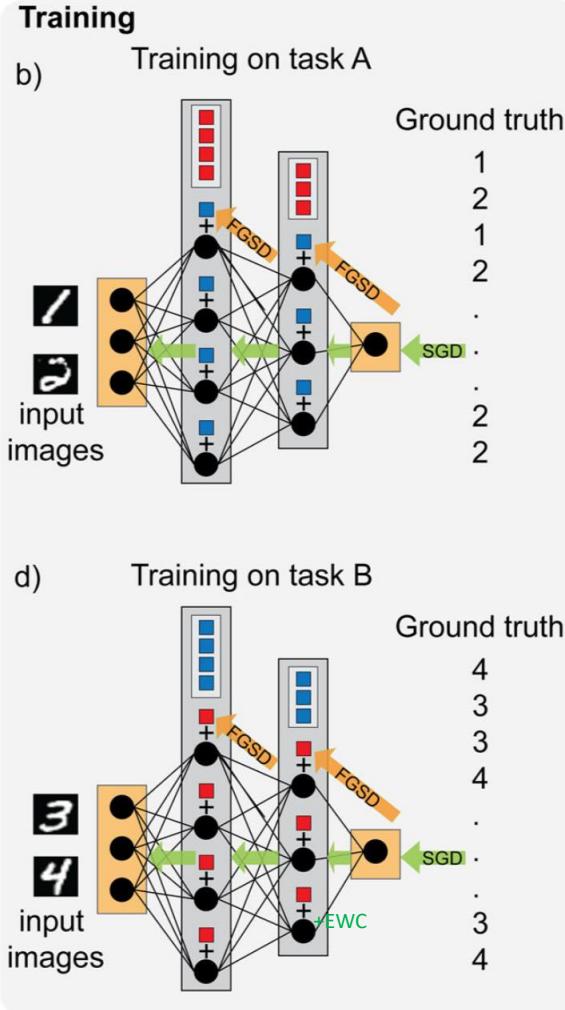
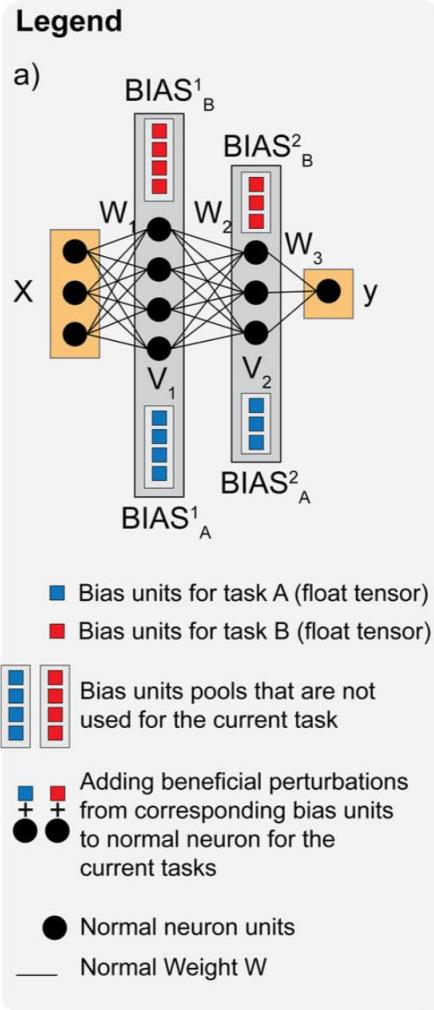
New proposed approach,  
Wen, Rios, Ge & Itti, 2021

# Overview of the approach

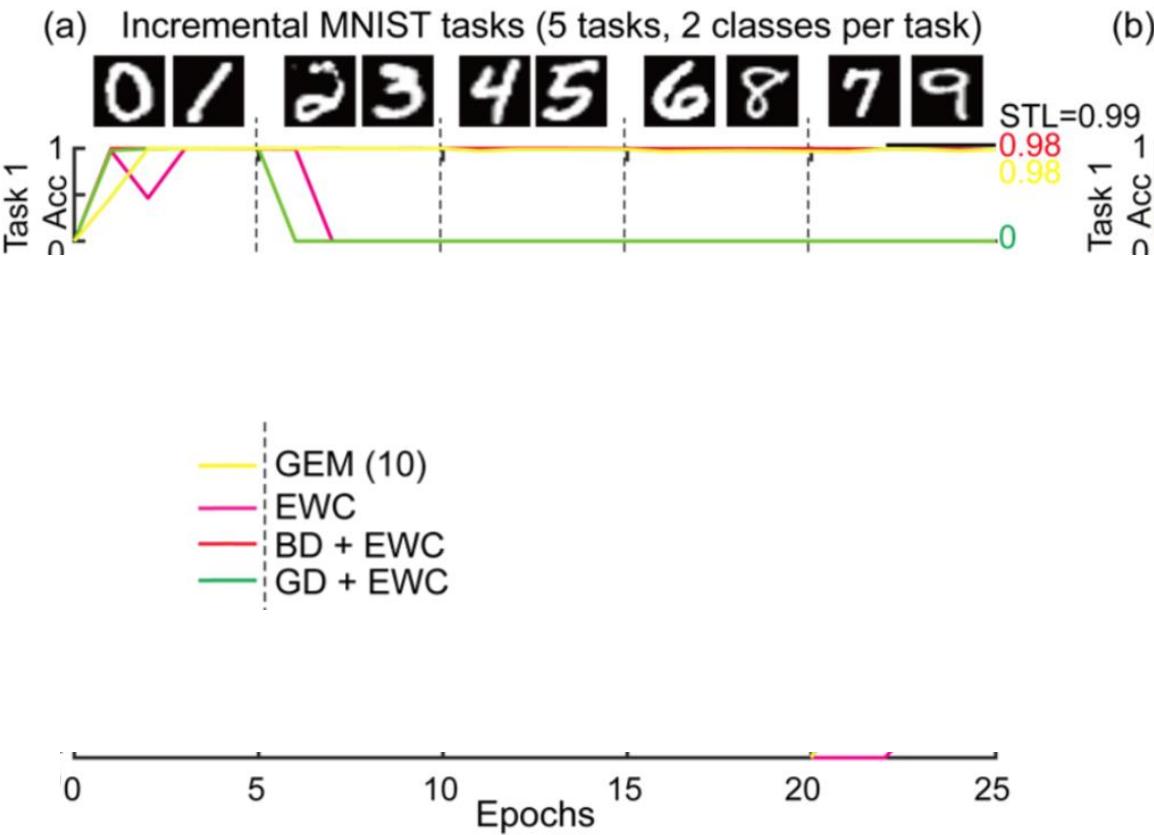
- Traditional learning e.g, SGD
- Continual learning e.g, EWC or PSP
- Beneficial perturbations for task A ( $BP_A$ )
- Beneficial perturbations for task B ( $BP_B$ )
- task A
- task B



# Overview of the approach

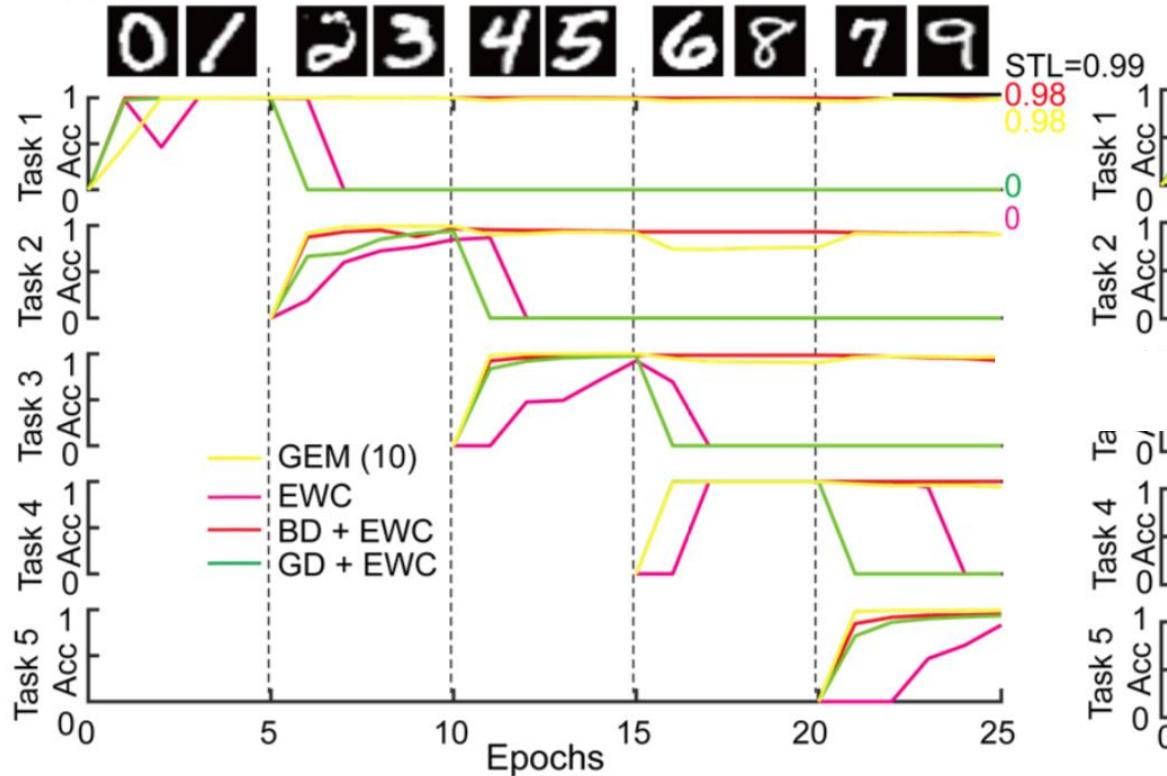


# Results

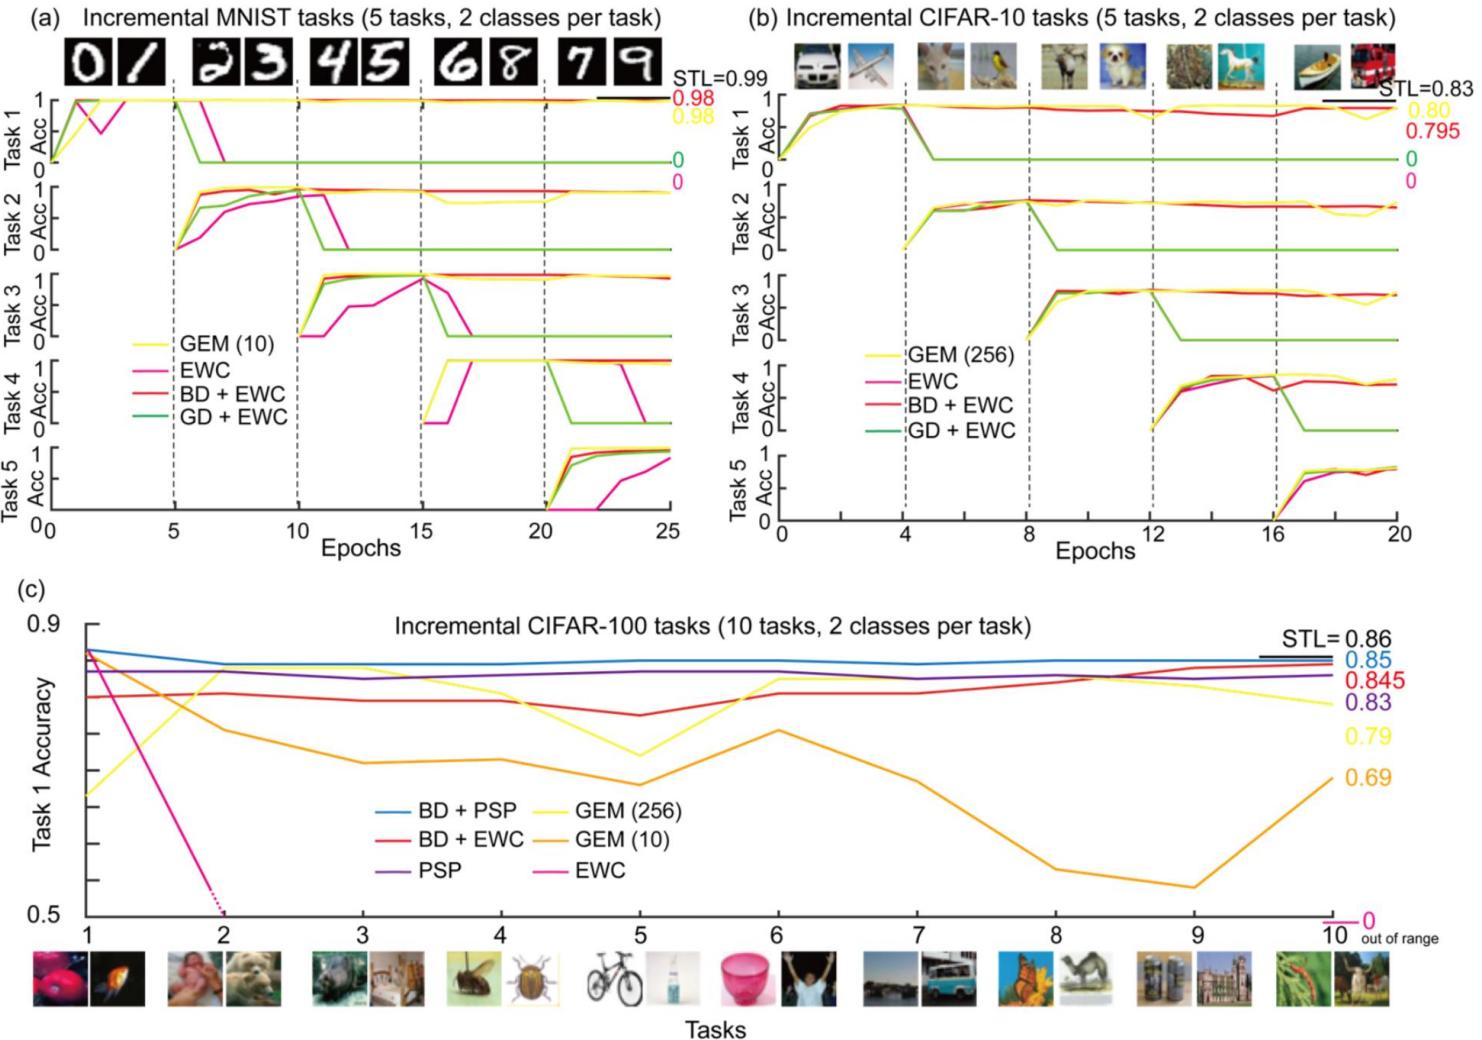


# Results

(a) Incremental MNIST tasks (5 tasks, 2 cla



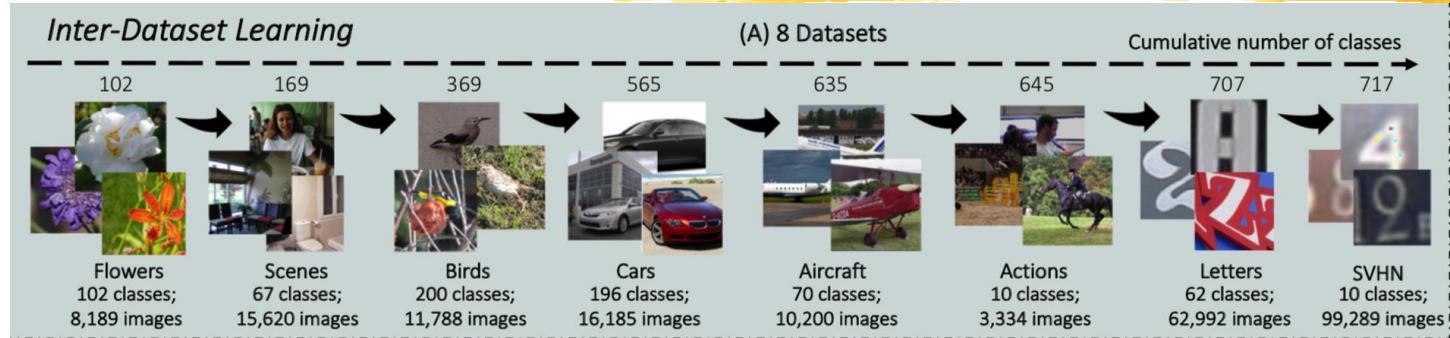
# Results



# Results

DATASET	METHOD	TASK I PERFORMANCE AFTER TRAINING ALL SEQUENTIAL TASKS	ADDITIONAL MEMORY STORAGE COSTS PER TASK (BYTES)
INCREMENTAL MNIST (5 TASKS, 2 CLASSES PER TASK)	GEM(10)	0.980	47,040
	BD+EWC	<b>0.980</b>	<b>4,808</b>
INCREMENTAL CIFAR-10 (5 TASKS, 2 CLASSES PER TASK)	GEM(256)	<b>0.800</b>	4,718,592
	GEM(150)	0.698	2,764,800
	BD+EWC	0.795	<b>4,808</b>
INCREMENTAL CIFAR-100 (10 TASKS, 2 CLASSES PER TASK)	GEM(256)	0.790	4,718,592
	GEM(209)	0.775	3,852,288
	BD+PSP	<b>0.850</b>	20,776
	PSP	0.830	15,968
	BD+EWC	0.845	<b>4,808</b>

# Results

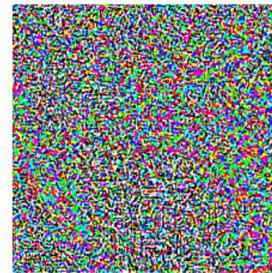


METHOD \ DATASET	FLOWER	SCENES	BIRDS	CARS	AIRCRAFT	ACTIONS	LETTERS	SVHN	AVERAGE
IMM [25]	67.44*	47.08*	42.28*	39.19*	18.93*	32.88*	46.35*	53.15*	43.41*
LwF [27]	65.68*	46.04*	42.47*	37.24*	30.75*	34.55*	50.31*	88.89*	49.49*
EWC [20]	-	-	-	-	-	-	-	-	50.00*
EBLL [45]	-	-	-	-	-	-	-	-	50.29*
SI [50]	74.58*	52.53*	48.65*	43.96*	32.19*	38.22*	47.7*	66.09*	50.49*
MAS [1]	76.34*	55.30*	49.03*	44.59*	33.48*	41.76*	51.93*	69.13*	52.69*
BD + EWC (OURS)	65.67	47.99	42.70	27.96	40.17	47.61	66.97	83.24	52.79
SLNID [2]	-	-	-	-	-	-	-	-	54.50*
PSP [5]	75.36	52.39	40.85	28.08	42.45	48.64	66.52	83.74	54.75
BD + PSP (OURS)	78.47	56.04	47.96	34.40	47.58	51.61	71.19	83.73	<b>58.87</b>

## Application: Adversarial robustness



$+ .007 \times$



=



“panda”

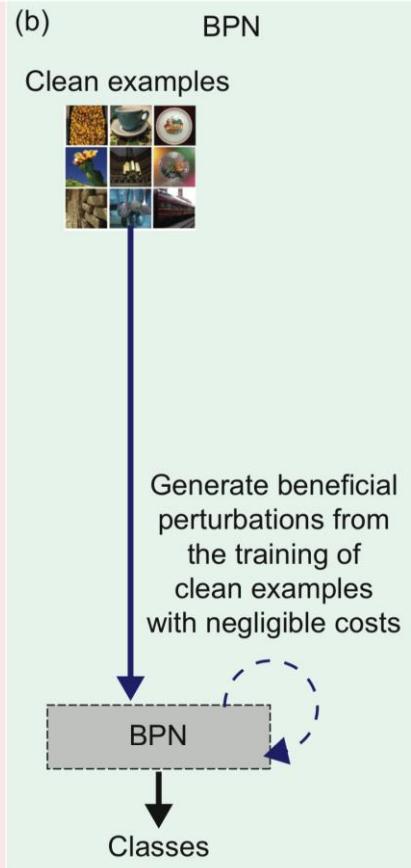
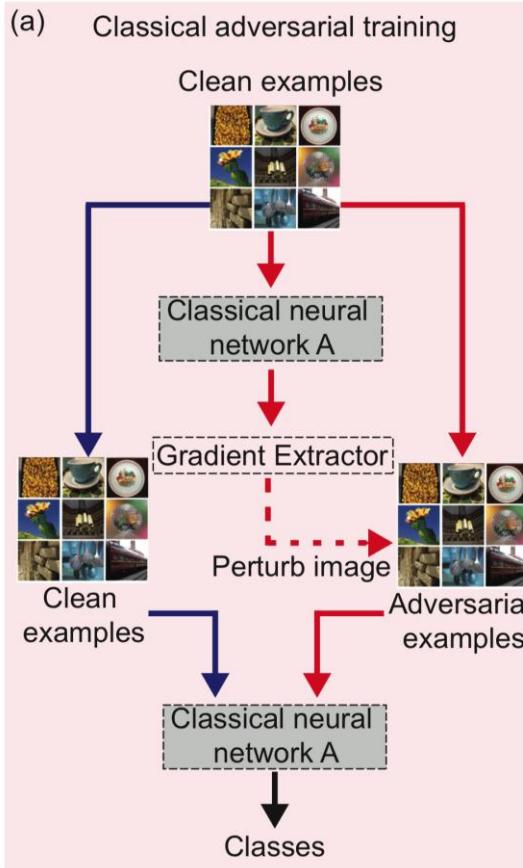
57.7% confidence

noise

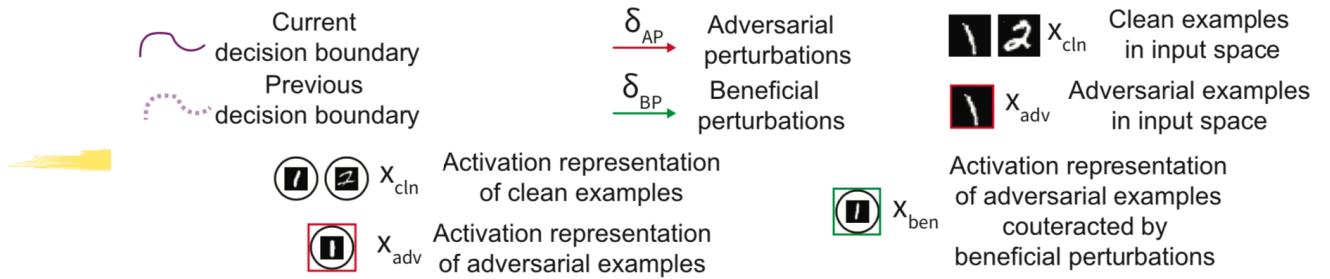
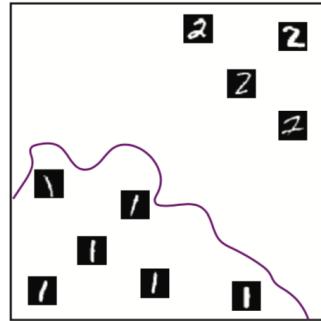
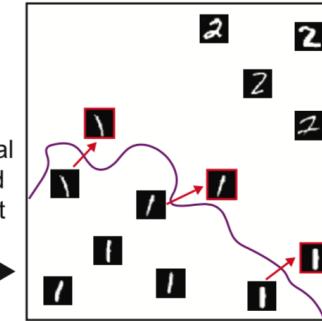
“gibbon”

99.3% confidence

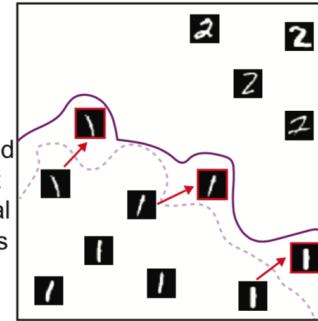
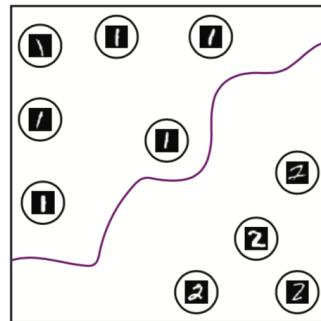
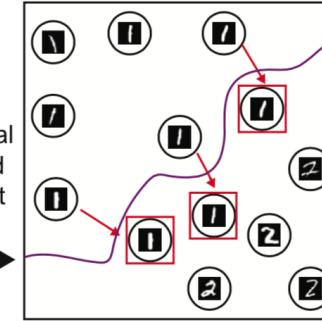
# Overview of the approach



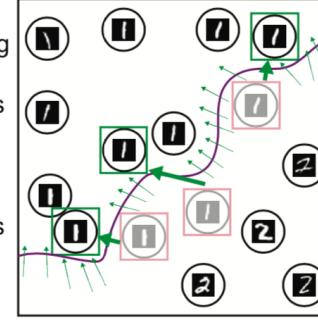
Wen & Itti, submitted

(a<sub>1</sub>) Input Space(a<sub>2</sub>) Input Space(a<sub>3</sub>) Input Space

Additional adversarial training needed to counteract adverversarial perturbations

(b<sub>1</sub>) Activation Space(b<sub>2</sub>) Activation Space(b<sub>3</sub>) Activation Space

Counteracting adversarial perturbations by already learned beneficial perturbations



## Results

network	Computation cost	Forward (FLOPS)	Backward (FLOPS)
Classical Network (ResNet-50)		51,112,224	51,112,225
BPN (ResNet-50)		51,112,224	51,115,321

## Results

Table 2. Training on clean examples for BPN and classical network (CN). Testing on clean examples (Cln Ex) and adversarial examples (Adv Ex) (generated by FGSM,  $\epsilon = 0.3$  for MNIST, FashionMNIST and TinyImageNet). CN does poorly on adversarial examples. While, BPN can successfully defend adversarial examples.

		Datasets & network structure	MNIST LeNet	FasMNIST LeNet	TinyImageNet ResNet-18
Cln Ex	BPN	<b>99.17</b>	<b>89.53</b>	57.55	
	CN	99.01	89.17	<b>64.30</b>	
Adv Ex	BPN	<b>98.88</b>	<b>54.07</b>	<b>53.29</b>	
	CN	18.08	11.87	1.45	

# Results

Table 5. Training on clean examples of MNIST and TinyImageNet for BPN and classical network (CN). Testing on adversarial examples generated by a variety of adversarial attack methods. BPN can successfully defend those adversarial examples.

Dataset & network		Attacks		PGD Linf	FGSM	Basic Iterative Attack L2	Aka Basic Iterative Attack
MNIST	BPN	<b>95.41</b>	<b>98.35</b>	<b>98.52</b>	<b>11.35</b>		
	CN	2.18	17.53	97.24	9.74		
TinyImageNet	BPN	<b>44.37</b>	<b>52.39</b>	<b>16.23</b>	<b>5.9</b>		
	CN	0.00	1.29	15.11	0.5		

# CLR: Channel-wise Lightweight Reprogramming for Continual Learning



Yunhao Ge



Yuecheng Li\*



Shuo Ni\*



Jiaping Zhao



Ming-Hsuan Yang



Laurent Itti

\* co-second author



Google Research

Paper



Code



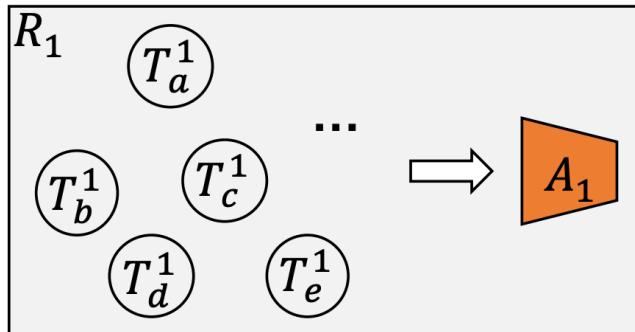
SKILL-102 dataset



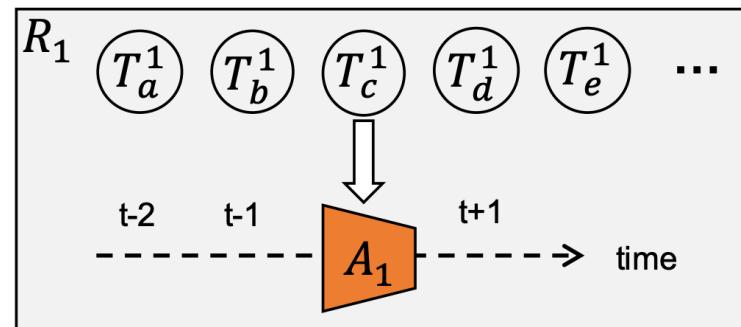
ICCV23  
PARIS

# Background: Continual Learning and Catastrophic Forgetting

## a) Multi-task Learning



## b) Sequential Lifelong Learning



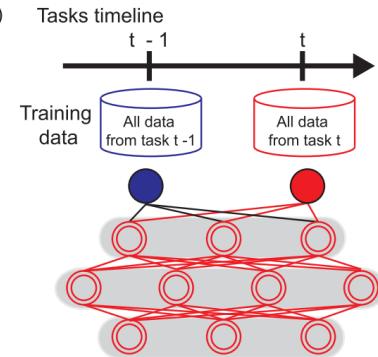
Ge, Yunhao, Yuecheng Li, Di Wu, Ao Xu, Adam M. Jones, Amanda Sofie Rios, Iordanis Fostiropoulos et al. "Lightweight Learner for Shared Knowledge Lifelong Learning." TMLR (2023)

Continual (Lifelong) learning aims to emulate the human ability to continually accumulate knowledge over sequential tasks.

# Background: 3 main types of Continual Learning methods

## Regularization methods

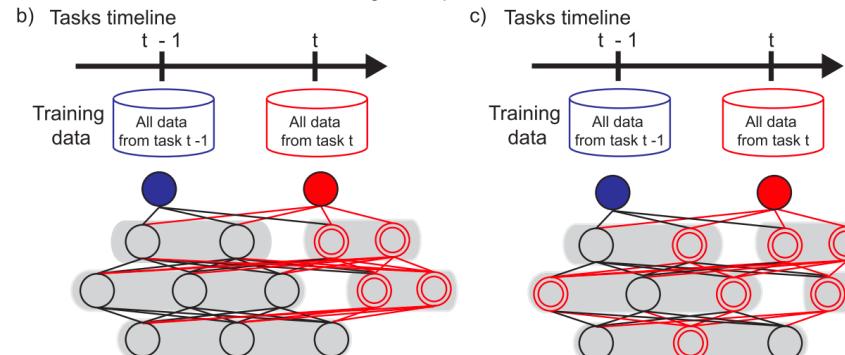
Type 1: constrain the network weights for previous tasks



EWC, IMM, SI, MAS,  
LwF, LFL, DMC

## Dynamic network methods

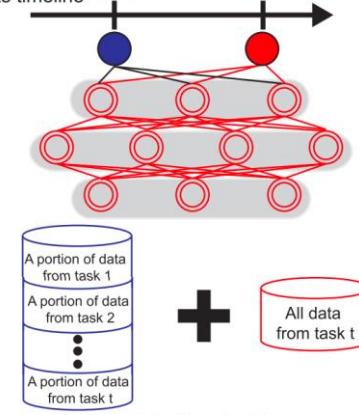
Type 2: expand the network capacity for new task and constrain the network weights for previous tasks



SUPSUP, PSP, CCLL, EFTs, BPN

## Replay methods

d) Type 3: store a subset of original dataset  
Tasks timeline t - 1      t



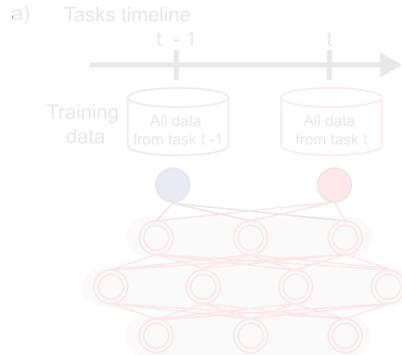
Train task t with episodic memory

iCaRL, ER GSS, AGEM,  
AGEM-R, DER, DERPP

# Background: 3 main types of Continual Learning methods

## Regularization methods

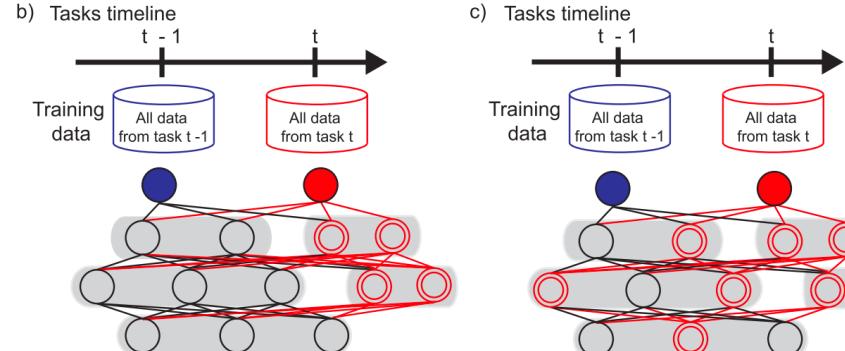
Type 1: constrain the network weights for previous tasks



EWC, IMM, SI, MAS,  
LwF, LFL, DMC

## Dynamic network methods

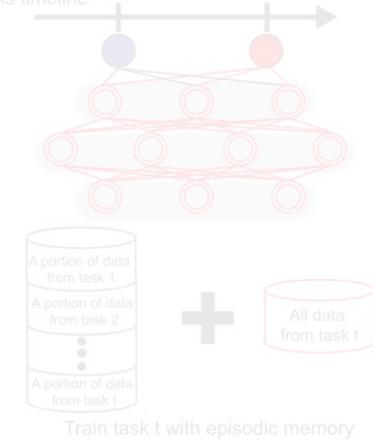
Type 2: expand the network capacity for new task and constrain the network weights for previous tasks



SUPSUP, PSP, CCLL,EFTs, BPN

## Replay methods

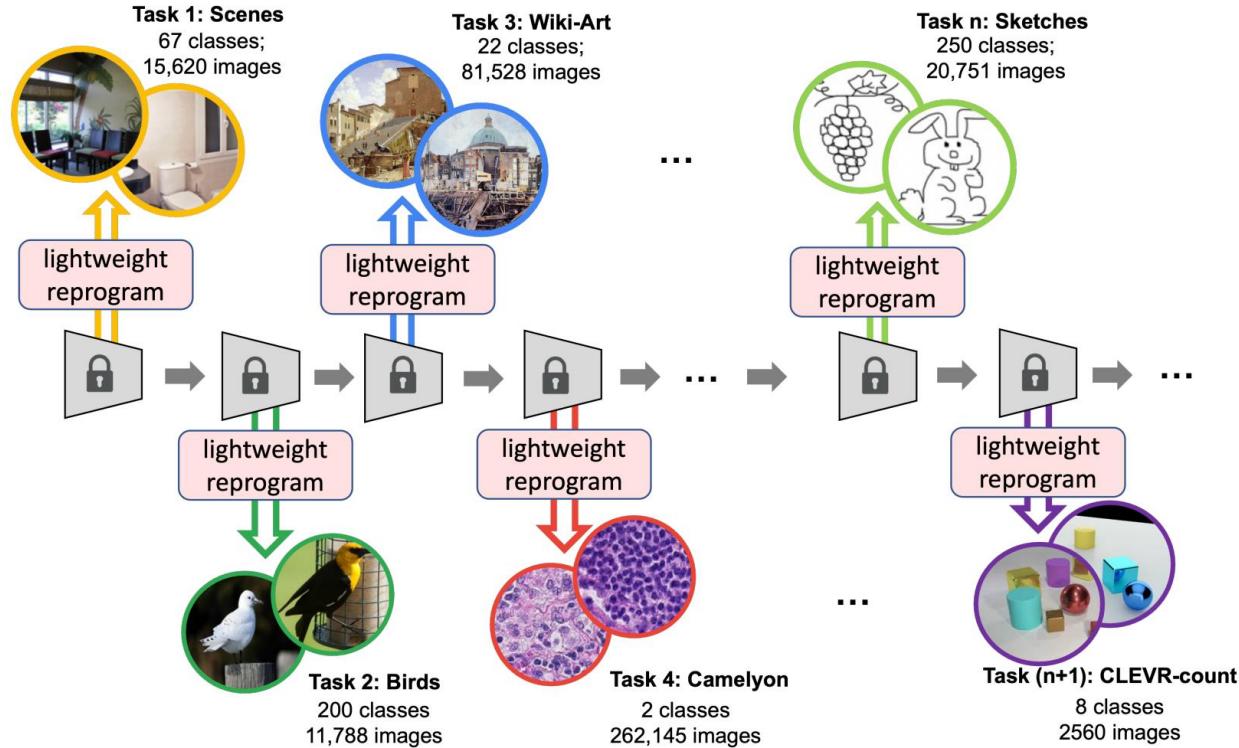
d) Type 3: store a subset of original dataset  
Tasks timeline t - 1 → t



Train task t with episodic memory

iCaRL, ER GSS, AGEM,  
AGEM-R, DER, DERPP

# CLR: Channel-wise Lightweight Reprogramming



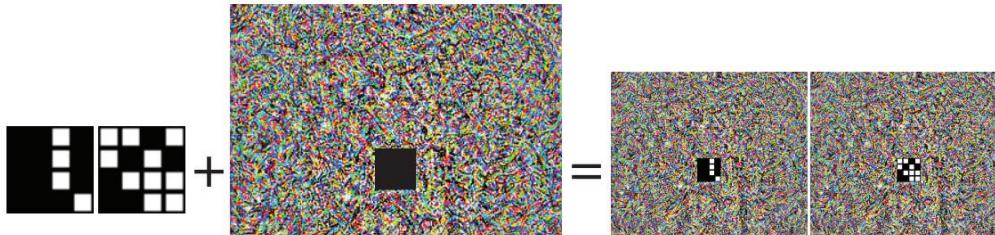
# Motivation 1: Adversarial Reprogramming

*Reuse instead of re-learn*

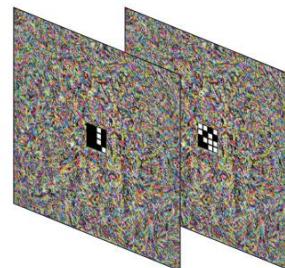
(a) counting ImageNet

$y_{adv}$	$y$
1 square	tench
2 squares	goldfish
3 squares	white shark
4 squares	tiger shark
5 squares	hammerhead
6 squares	electric ray
7 squares	stingray
8 squares	cock
9 squares	hen
10 squares	ostrich

(b) Adversarial Program



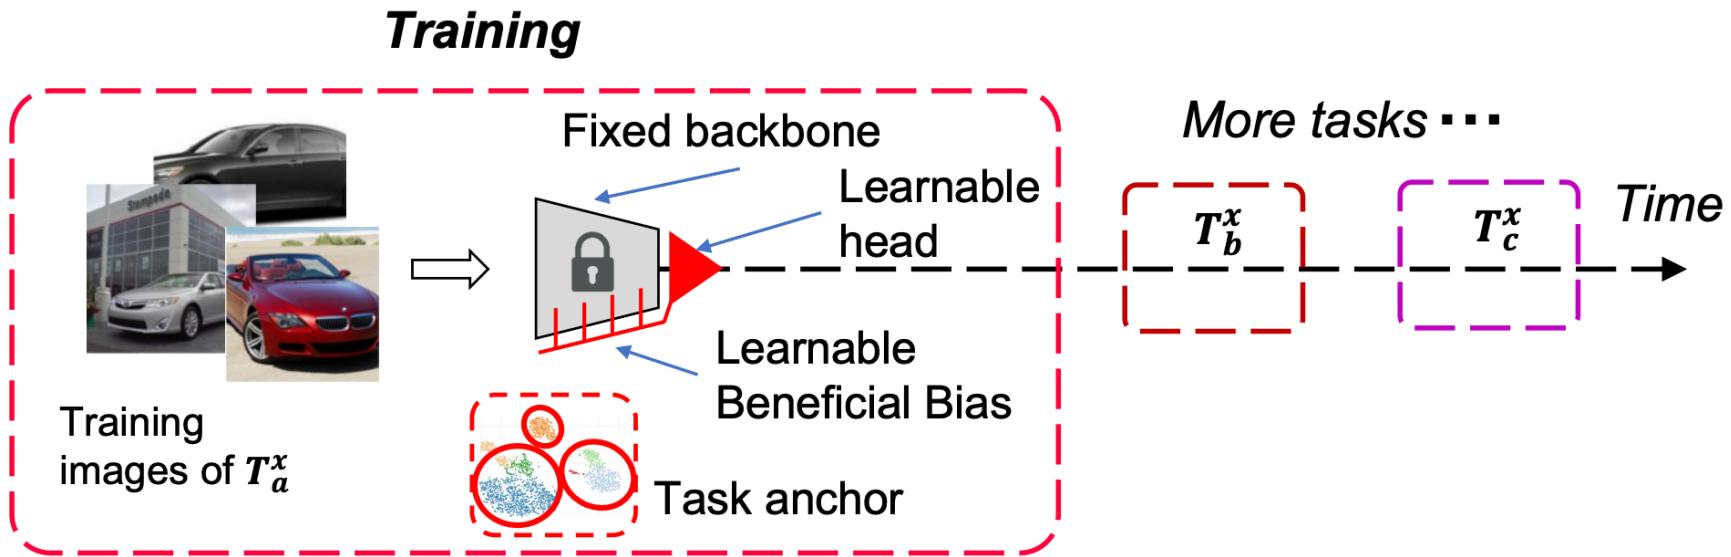
(c)



tiger shark, ostrich  
≡  
4 squares, 10 squares

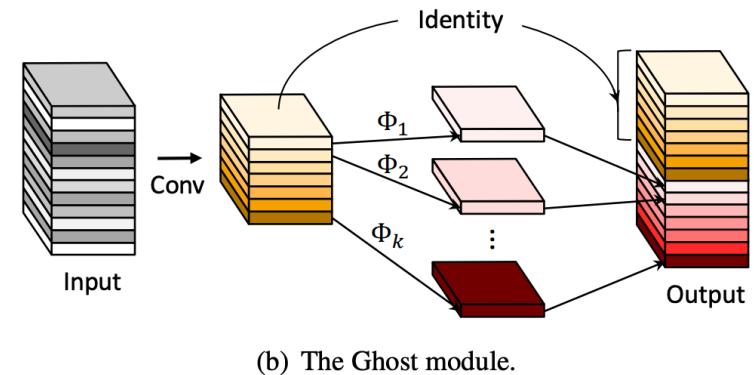
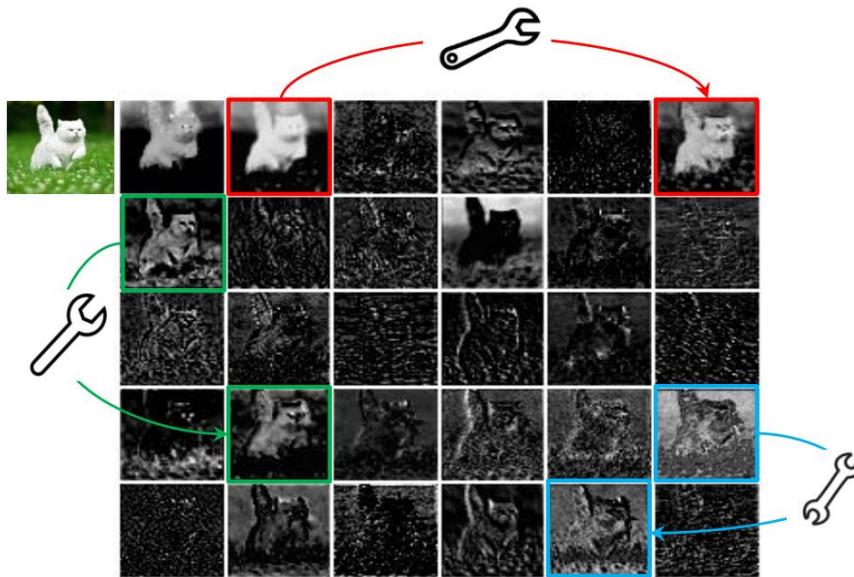
# Motivation 2: Light-weight Lifelong Learner

***Lightweight parameters could shift model distribution***

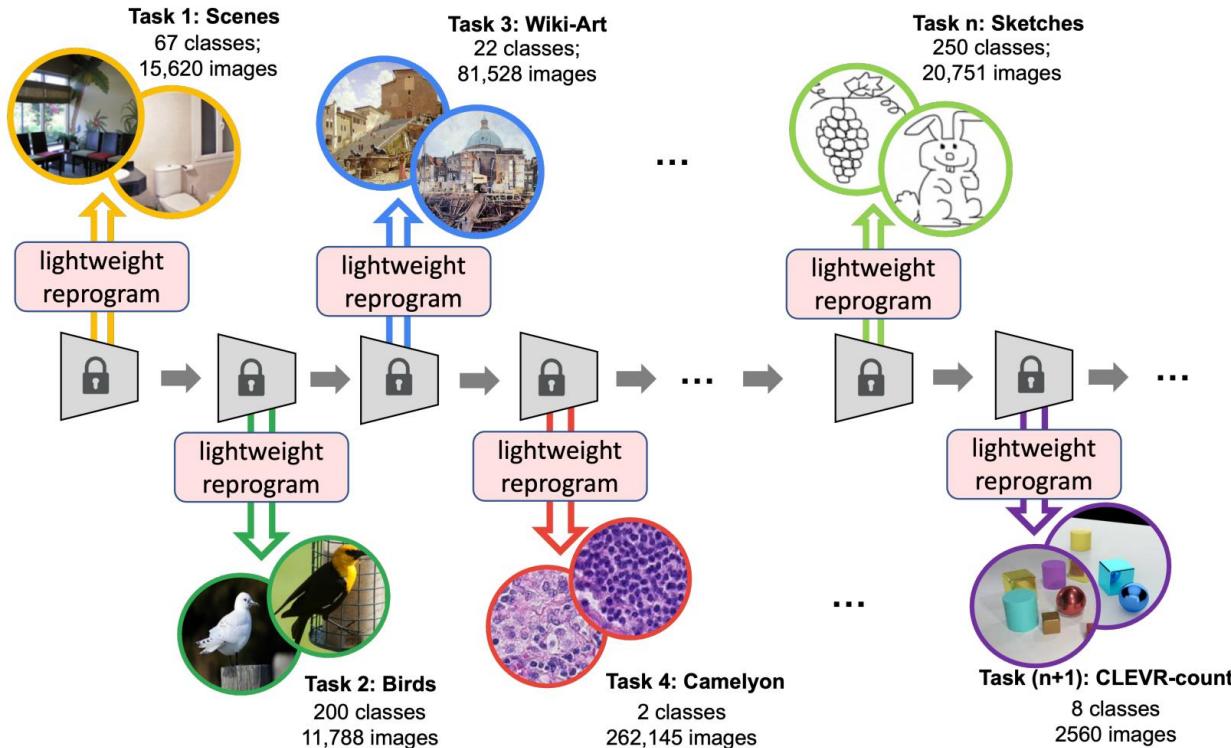


# Motivation 3: GhostNet

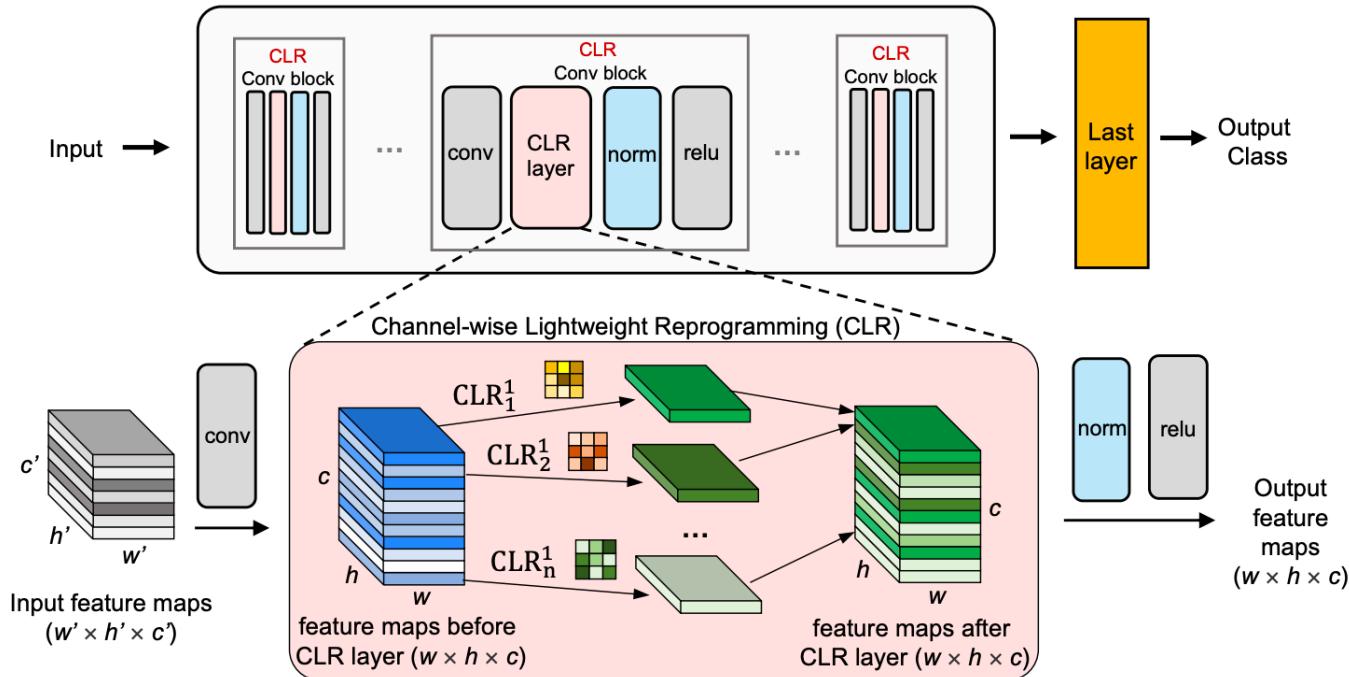
*Channel-wise transformations may link two different kernels*



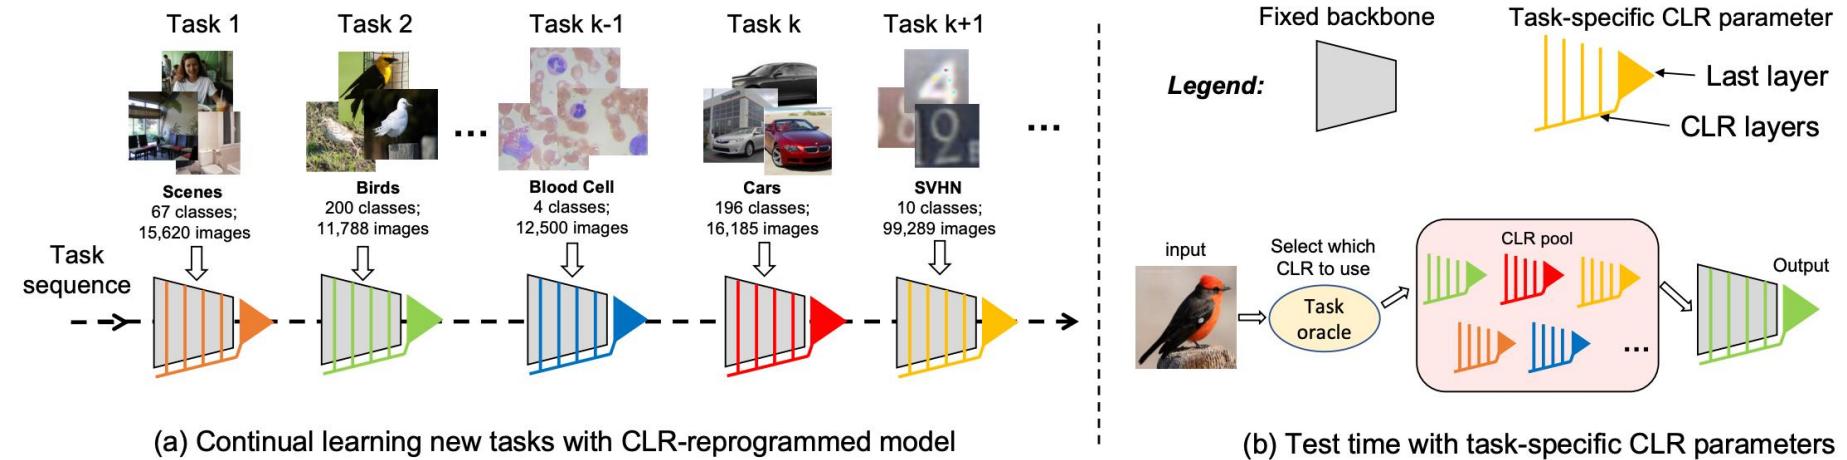
# CLR: Channel-wise Lightweight Reprogramming



# CLR: Channel-wise Lightweight Reprogramming

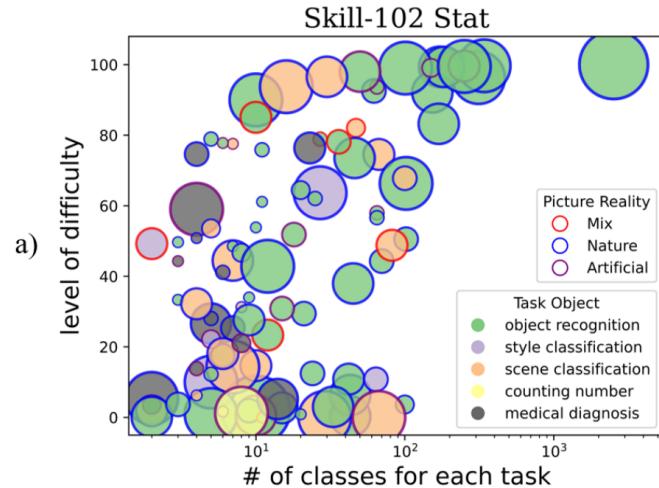


# CLR: Channel-wise Lightweight Reprogramming

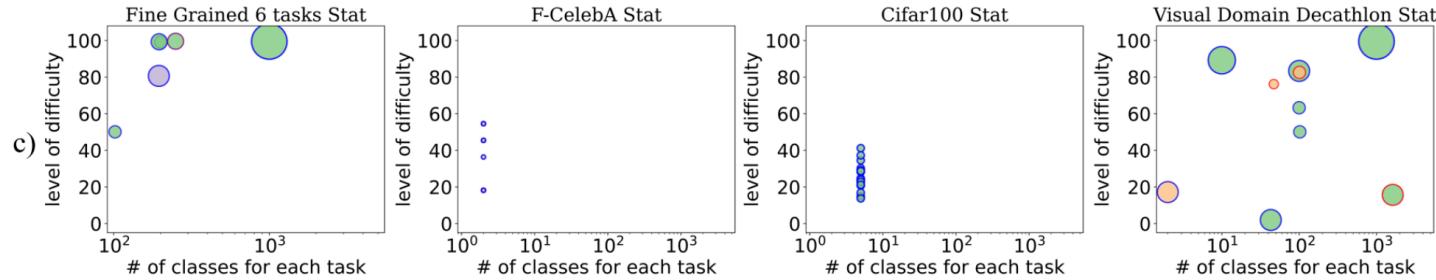


# Experiments: SKILL-102 continual Learning dataset

SKILL-102 dataset



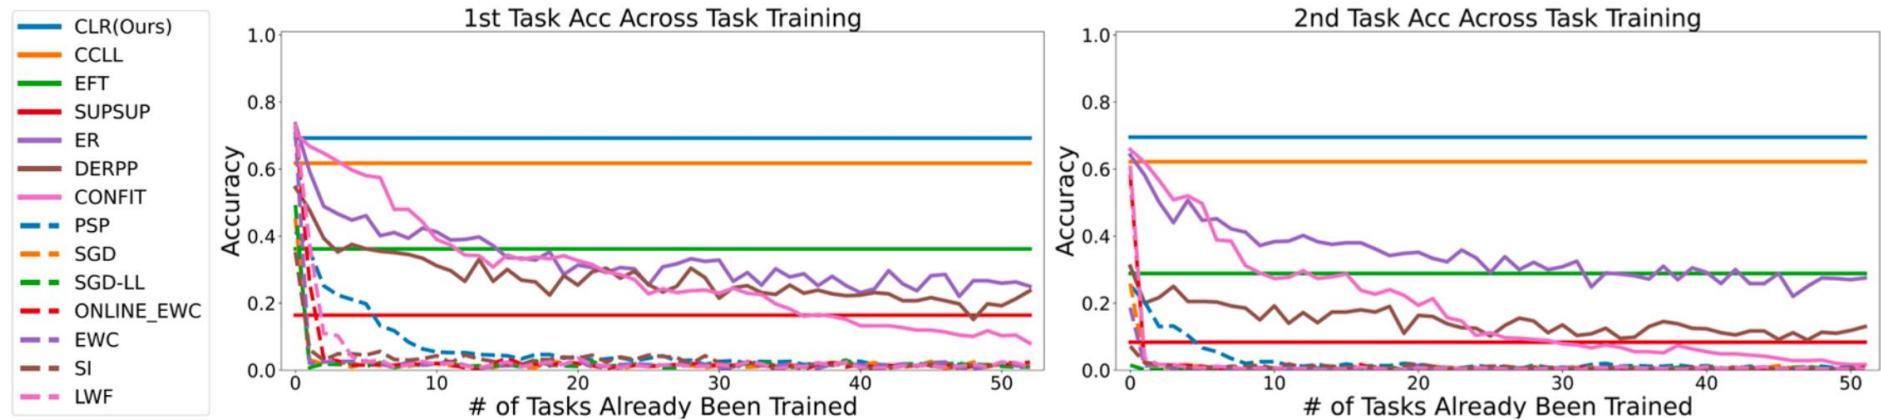
Comparison	SKILL-102	Visual domain decathlon	Cifar100	F-CelebA	Fine grained 6 tasks
# Tasks	<b>102</b>	10	20	10	6
# Classes	<b>7,059</b>	3,128	100	20	1,943
# Images	<b>5,162,149</b>	1,659,142	60,000	1,180	1,440,086
Max( $\Delta$ accuracy/difficulty)	<b>99.60%</b>	97.56%	27.45%	36.36%	49.50%
# Different classification target	<b>5</b>	2	1	1	2
Mix image style (nature/artifact)	✓	✓	✓	✗	✓
Mix super/fine-class classification	✓	✓	✗	✗	✓



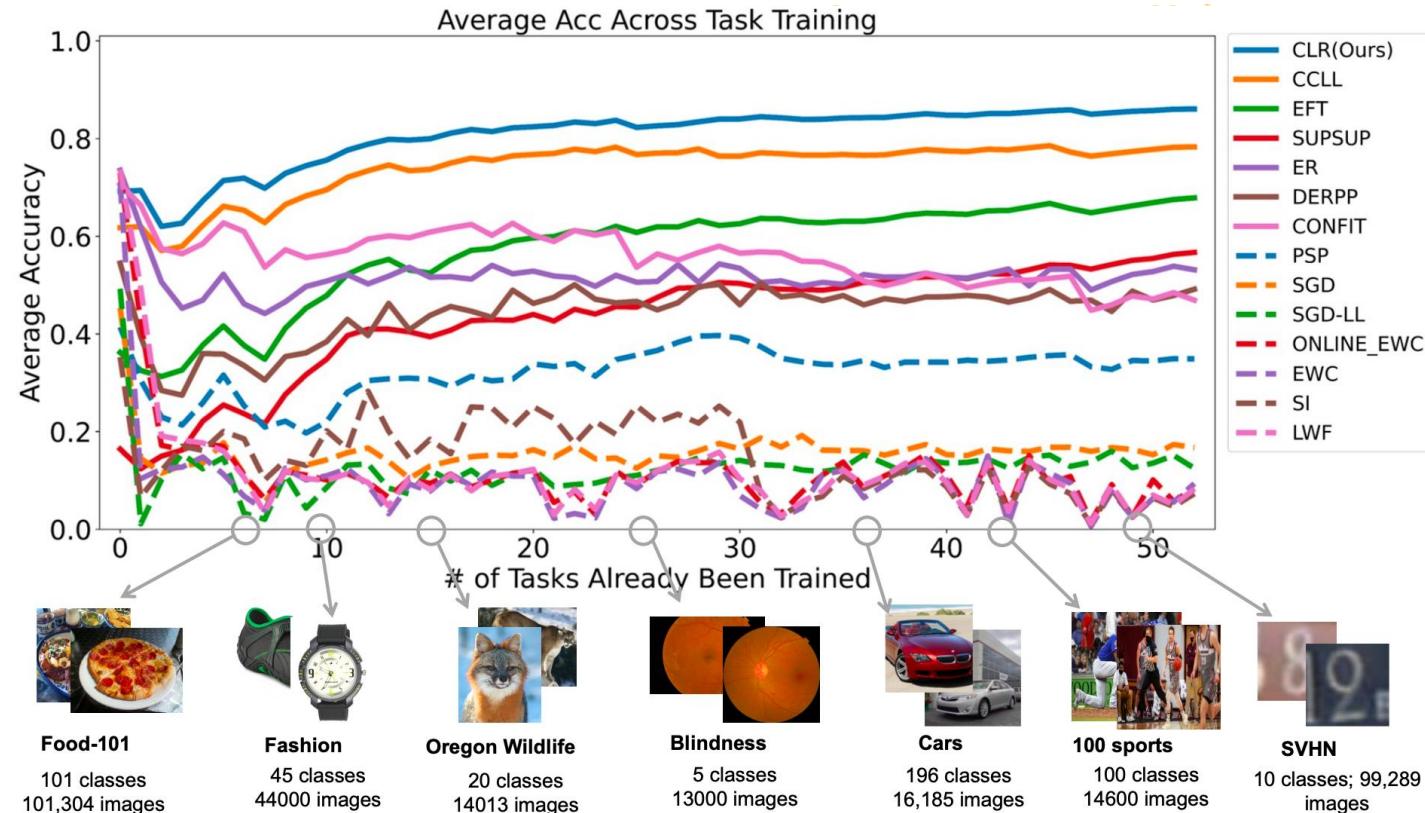
# Experiments: SKILL-102 continual Learning dataset

Comparison	53-dataset (ours)	8-dataset [1, 2]	ImageNet [47]	Fine grained 6 tasks [47] [35]	Cifar100 [34]	F-CelebA [43]
# tasks	<b>53</b>	8	20	6	20	10
# Classes	1,584	738	1000	<b>1943</b>	100	20
# Images	<b>1,811,028</b>	166,360	1,300,000	1,440,086	60,000	1,189
# different classification target	5	1	1	2	1	1
Mix image style (nature/artifact)	✓	✓	✗	✓	✗	✗
Mix super/fine-class classification	✓	✓	✓	✗	✗	✗

# Experiments: First task performance during continual learning



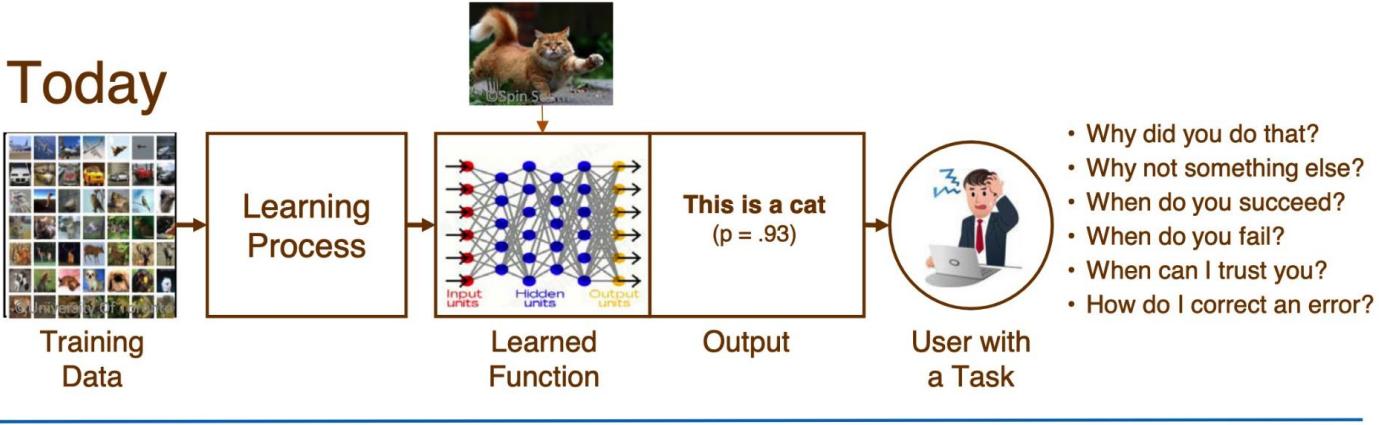
# Experiments: Average accuracy on all tasks learned so far



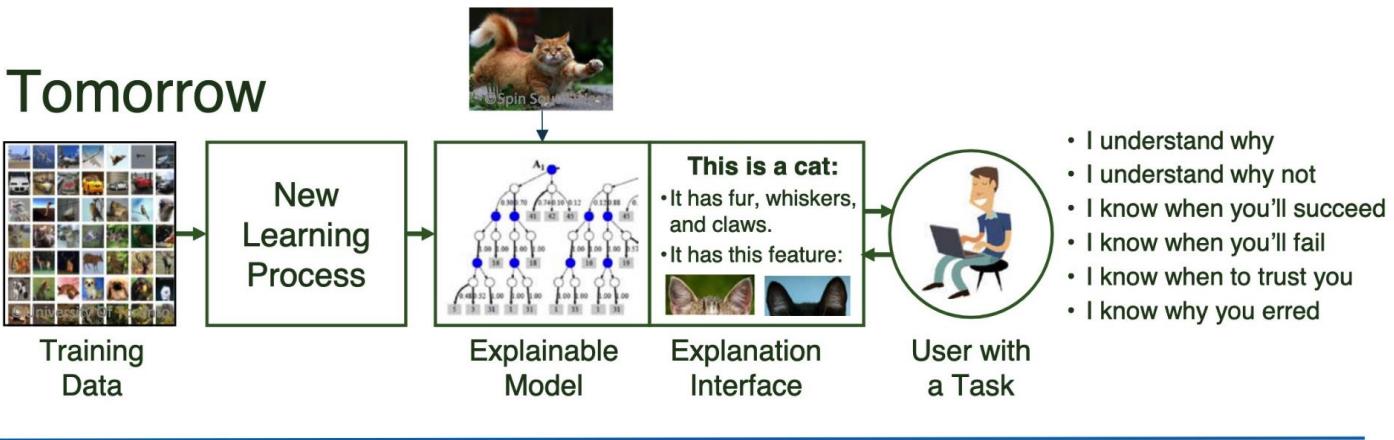
# Experiments: Parameter and computation cost

Method	Extra parameter to add 1 new task	Computation cost	Average Acc (53-datasets)
SGD	<b>0 %</b>	1	16.71 %
PSP [12]	5.02%	0.828	34.91 %
EWC [32]	0%	1.160	9.36 %
ONLINE EWC [50]	0%	1.011	9.47* %
SGD-LL	<b>0 %</b>	<b>0.333</b>	12.49 %
ER [45]	189.14%	3.99	53.13 %
SI [60]	0%	1.680	7.28%
LwF [36]	0%	1.333	8.23%
SUPSUP [58]	3.06%	1.334	56.69 %
EFT [55]	3.17%	1.078	67.8 %
CCLL [53]	0.62%	1.006	78.3 %
CLR (Ours)	0.59%	1.003	<b>86.05 %</b>

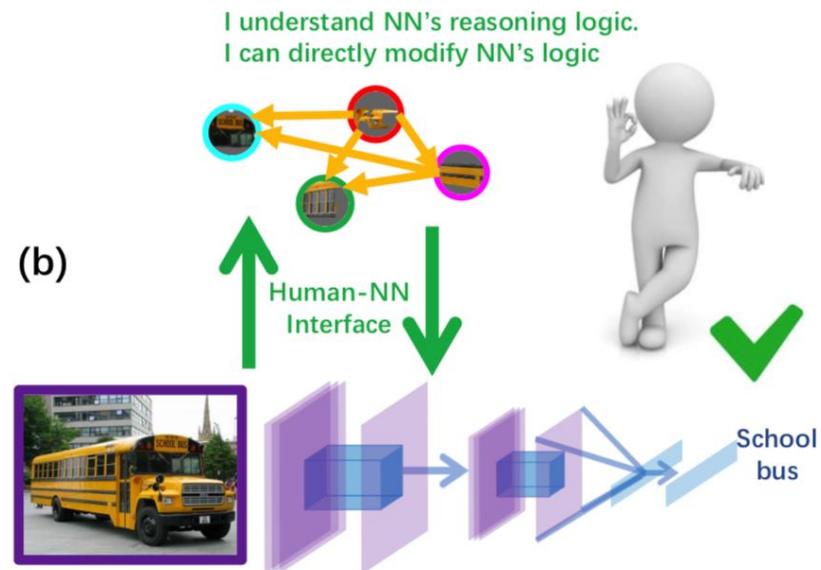
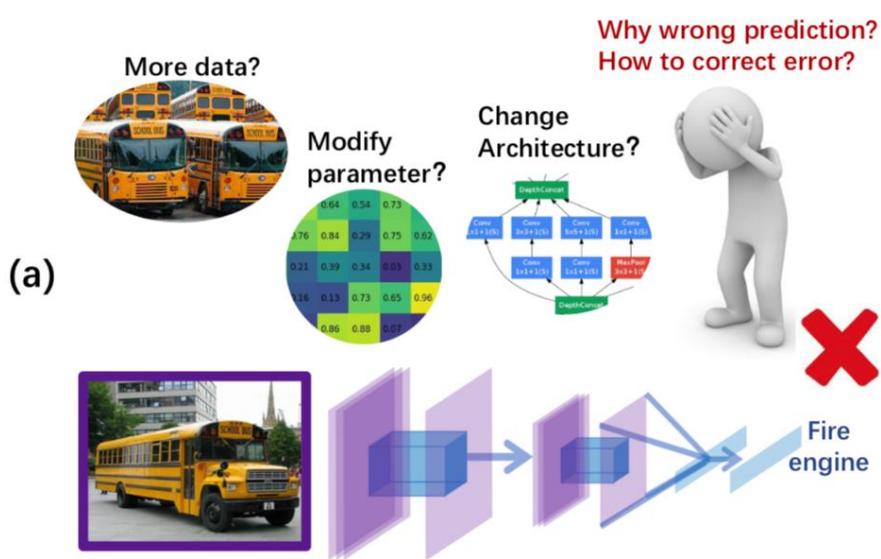
## Today



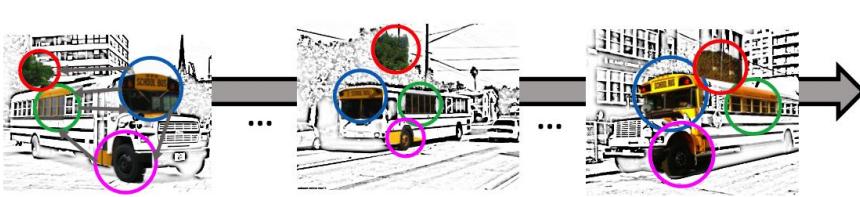
## Tomorrow



# Robustness through explainability

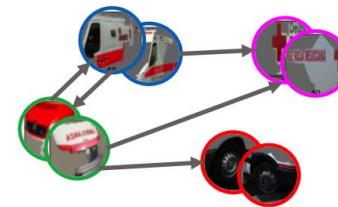
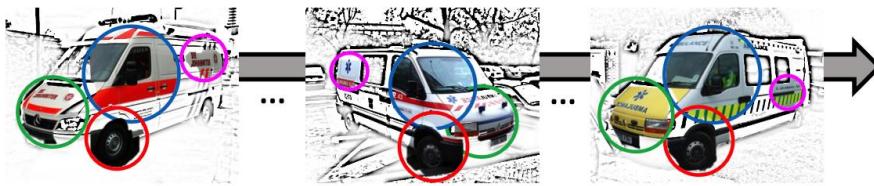


School bus  
image samples



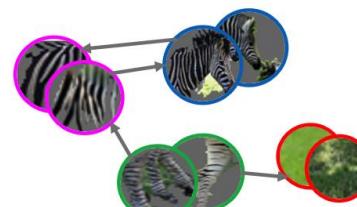
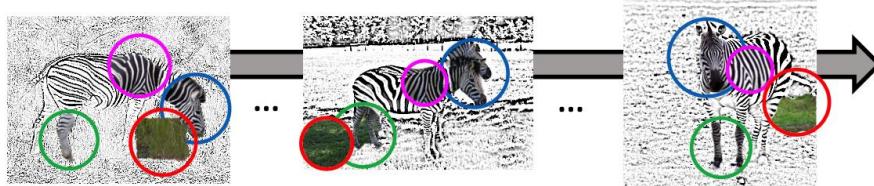
School bus  
c-SCG

Ambulance  
image samples



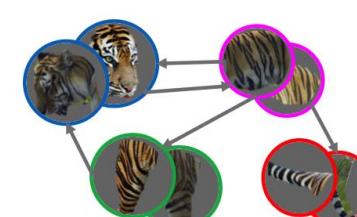
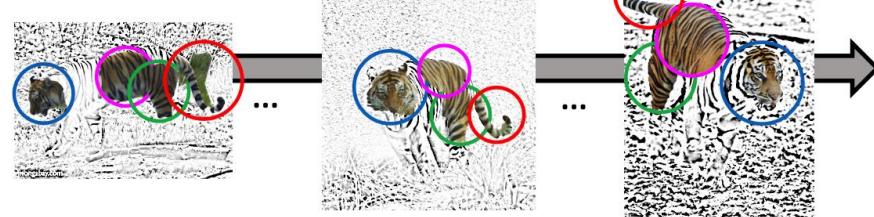
Ambulance  
c-SCG

Zebra  
image samples



Zebra  
c-SCG

Tiger  
image samples

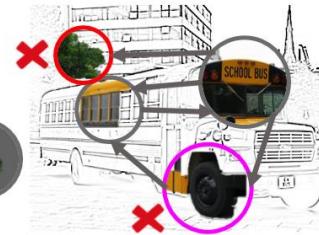
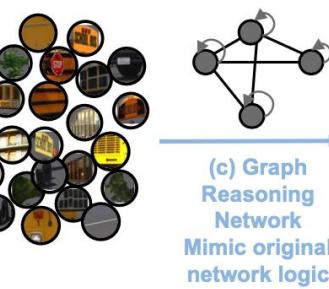
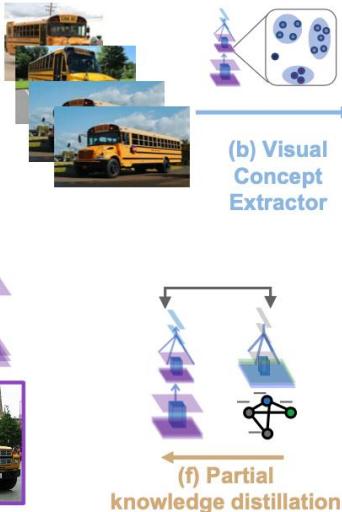
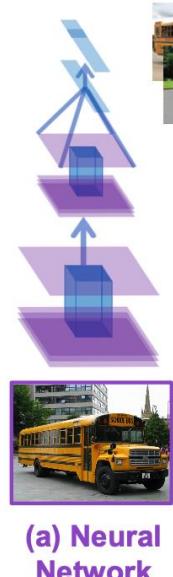


Tiger  
c-SCG

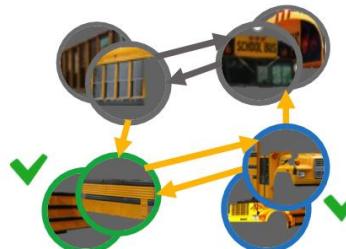
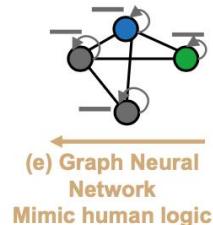
# Robustness through explainability

## Human-Network Interface

Network shows reasoning logic to human

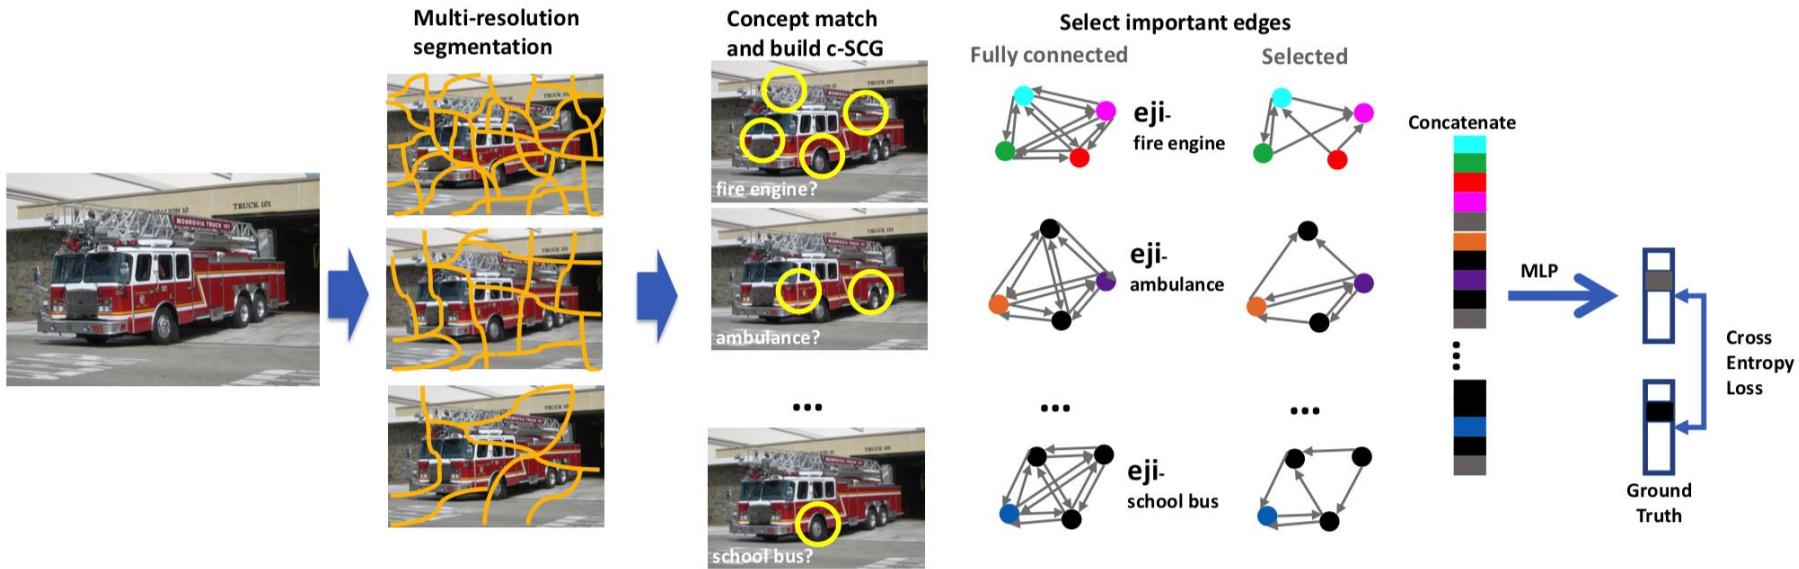


(d) Human Intelligence



Human change Network's reasoning logic

# Extracting visual concept graphs



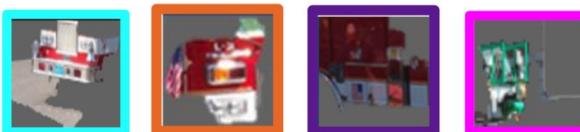
## Example: human modifies nodes



Human  
modify

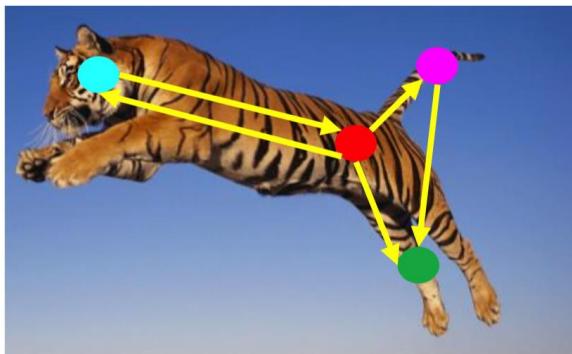


Original fire engine c-SCG

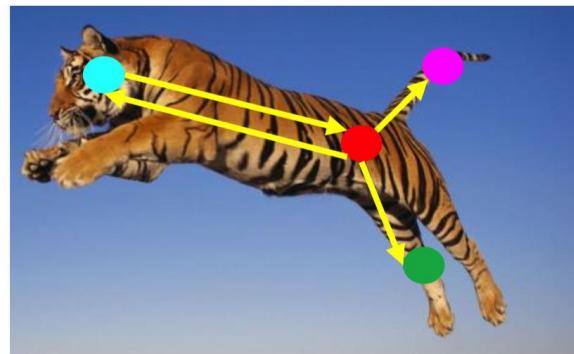


New Fire engine c-SCG

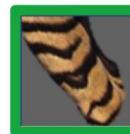
## Example: human modifies edges



Human  
modify

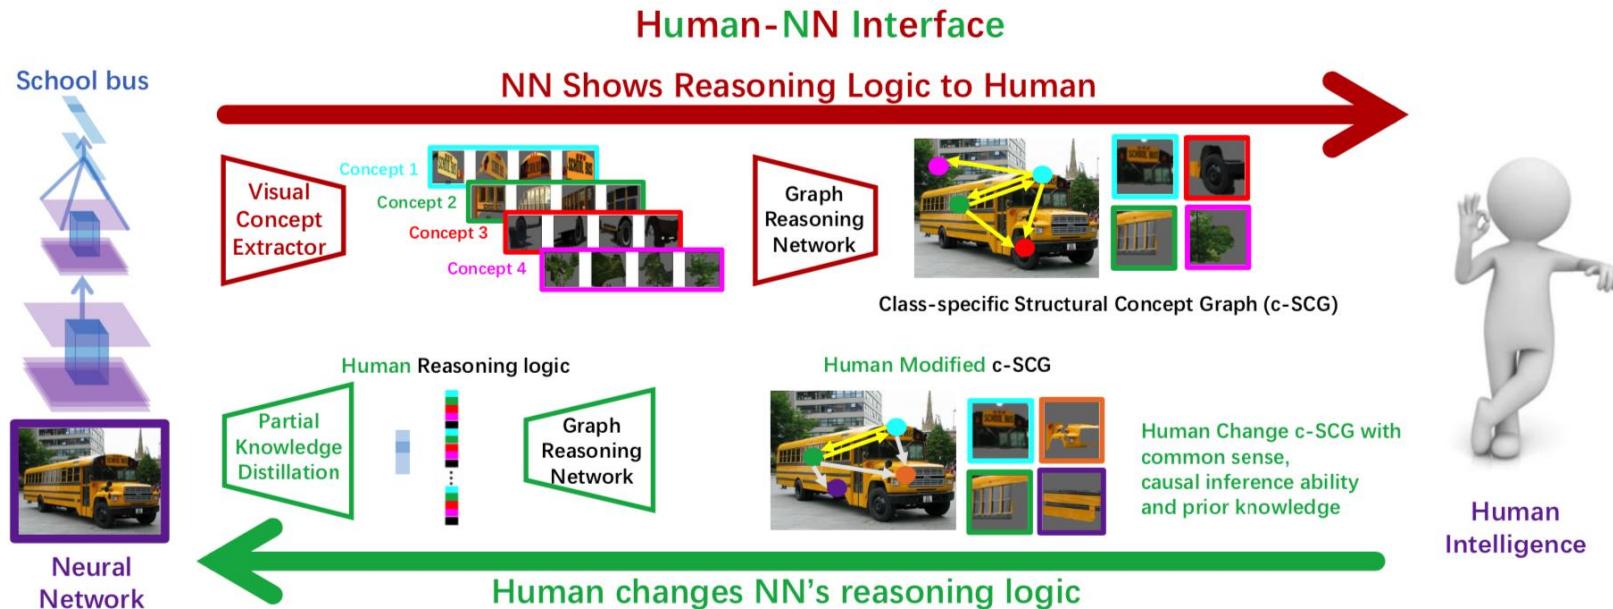


Original tiger c-SCG

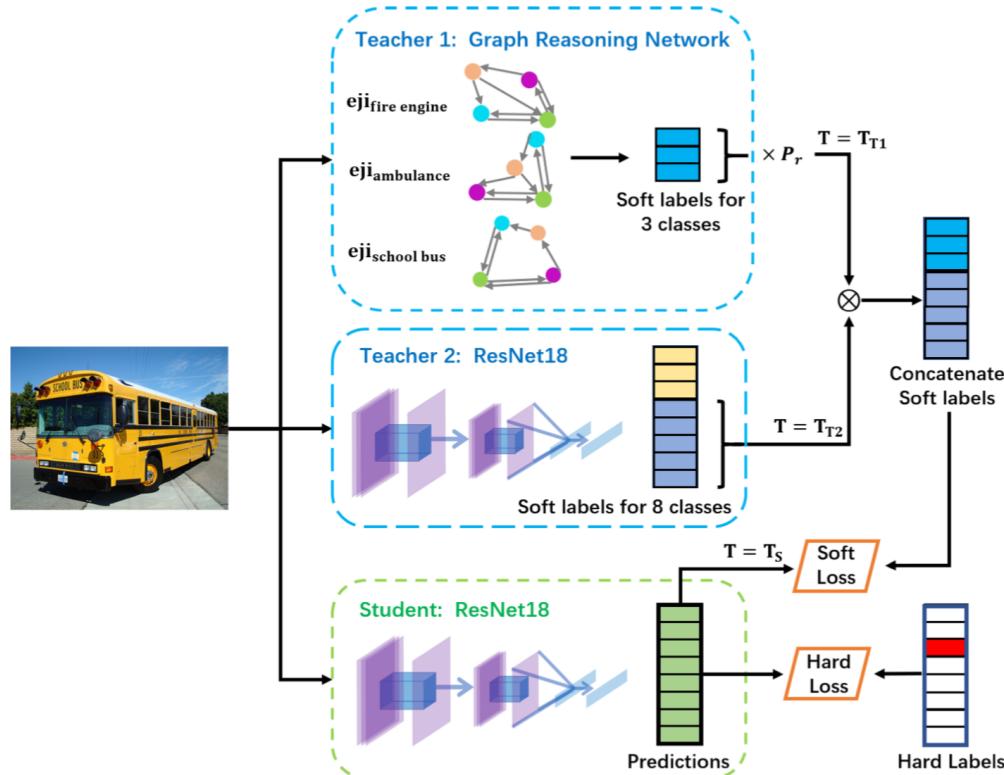


New tiger c-SCG

# Robustness through explainability



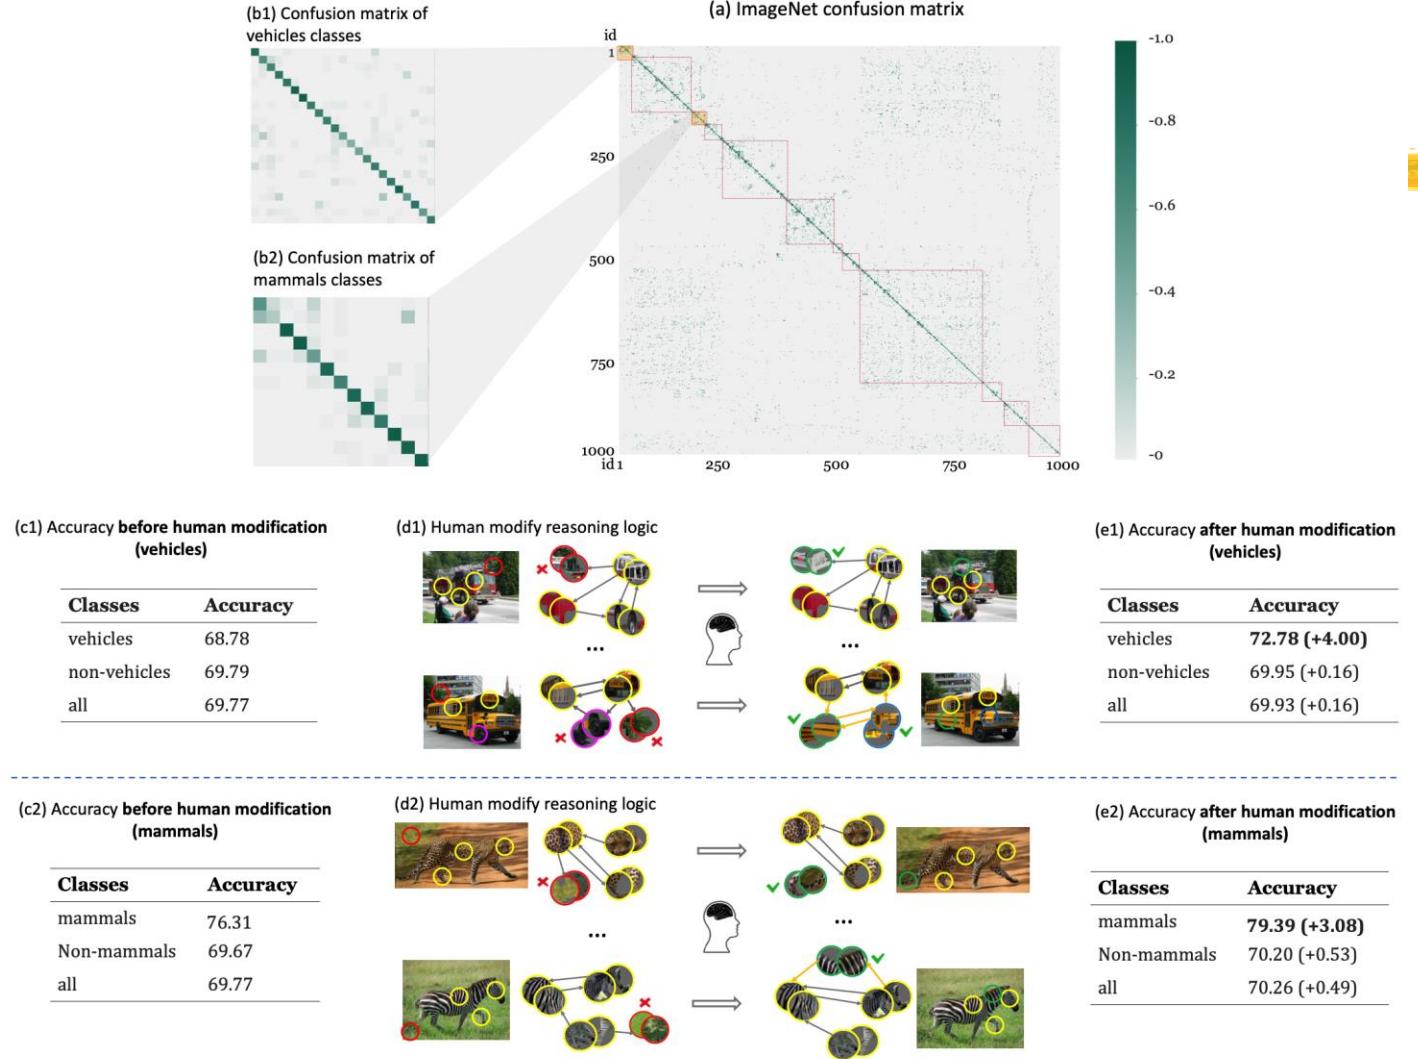
# Partial knowledge distillation



# Results

Table 1: Human improves a network's performance with HNI: experiments on six different image classification tasks (performance is tabulated as percent correct classification).

Datasets	# images	Classes	# classes	original performance	performance with HNI
Cats	2382	modified classes	2	88.33	91.64 (+3.31)
		all classes	12	93.06	93.61 (+0.55)
Cars	8144	modified classes	2	83.33	91.67 (+8.34)
		all classes	10	86.33	88.33 (+2.00)
Monkeys	1642	modified classes	3	78.75	85.00 (+6.25)
		all classes	10	90.00	93.61 (+3.61)
Flowers	22267	modified classes	3	78.94	85.76 (+6.82)
		all classes	10	82.37	85.57 (+3.2)
Fashion	16186	modified classes	3	61.71	67.14 (+5.43)
		all classes	6	73.79	74.37 (+0.58)
Buildings	5063	modified classes	3	47.78	57.78 (+10.00)
		all classes	17	53.73	54.51 (+0.78)



**Figure 3: Humans can improve a network's performance with HNI. We conduct large-scale experiments on the ImageNet dataset, which contains 1,000 real-world classes.**

(a) Confusion matrix of a 1,000-class original GoogleNet image classification network trained on ImageNet. Most of the errors are within each of 12 super-classes that correspond to groups of related classes (e.g., mammals, vehicles, birds, etc.). There are two main challenges: (1) How to correct the errors and improve accuracy within a super-class (local logic) with the help of human involvement? (2) How to maintain the performance of all other classes in the 1,000 classes? We show results of two large-scale experiments, for the superclasses of vehicles (23 classes, total 23,000 training images) and mammals (13 classes, 13,000 images), to show how one can use HNI to improve performance within a superclass without degrading performance of other classes. We first consider the super-class of vehicles (b1), with original accuracy over these 68.78% (c1). For each class, the network-to-human pass was used to show the reasoning logic of the original network as a c-SCG to a human operator (d1). The operator spotted and corrected any reasoning errors of the network. The human-to-network pass then distilled the human-modified logic back to the original network with the help of Graph neural network and partial knowledge distillation. Performance was improved on the vehicle classes, without degradation of non-vehicle classes (e1), demonstrating how humans could use their own knowledge to correct reasoning errors of the network and improve network accuracy. The same process is also shown for the superclass of mammals (b2, c2, d2, e2).

# Conclusions

---

- New Beneficial Biases help not only with lifelong learning but also with robustness to adversarial attacks
  
- Human/NN interaction through graph representations allows one to improve robustness and generalization by injecting human knowledge about natural vs spurious correlations in training data.

This work was supported by C-BRIC (one of six centers in JUMP, a Semiconductor Research Corporation (SRC) program sponsored by DARPA), DARPA (HR00112190134) and the Army Research Office (W911NF2020053). The authors affirm that the views expressed herein are solely their own, and do not represent the views of the United States government or any agency thereof.

# Some slides in this lecture are from

<https://courses.grainger.illinois.edu/ece448/sp2021/slides/lec25.pdf>

## CS440/ECE448 Lecture 25: AI Ethics



Modified by Mark Hasegawa-Johnson, 4/2021

Including slides by Svetlana Lazebnik, 10/2017

CC-BY-4.0: Copy at will, but cite the source

Image source: [https://www.britac.ac.uk/  
audio/machines-morality-and-future-medical-care](https://www.britac.ac.uk/audio/machines-morality-and-future-medical-care)

# Problem statement: Example 1



- Bickel, Hammel, and O'Connell, "Sex bias in graduate admissions: Data from Berkeley," Science 187(4175):398–404, 1975
- Women were being admitted to Berkeley at a far lower rate than men.

# Problem statement: Example 1



- Bickel, Hammel, and O'Connell, "Sex bias in graduate admissions: Data from Berkeley," Science 187(4175):398–404, 1975
- Women were being admitted to Berkeley at a far lower rate than men.
- Women were applying to departments with lower acceptance rates. Within each department, admission rates for men and women were the same.
- Is this fair?

# Problem statement: Example 2



A given job might discriminate based on:

- Upper body strength
- Habitual clothing
- Undergraduate major

All of these correlate with gender.

Is it fair to use them as a basis for employment?

## OpenAI faces complaint to FTC that seeks investigation and suspension of ChatGPT releases

The Center for AI and Digital Policy accuses OpenAI of violating a part of the FTC Act that prohibits unfair and deceptive business...

 Reuters

## Australian mayor readies world's first defamation lawsuit over ChatGPT content

A regional Australian mayor said he may sue OpenAI if it does not correct ChatGPT's false claims that he had served time in prison for...

5 days ago

 Times of India

## China's Baidu has sued Apple over fake versions of its ChatGPT rival Ernie bot

Chinese search engine giant Baidu has filed lawsuits against "relevant" app developers and Apple over fake copies of its ChatGPT rival Ernie...

2 days ago



# AI ethics, fairness, and bias



## Fairness

- Representation
- Algorithmic solutions
- Transparency as a solution to the fairness problem

## Accountability

- Physical safety – Data safety
- Economic safety

## Transparency

- Explainable AI
- Understanding the inner workings of a superintelligence

# AI ethics, fairness, and bias



## Fairness

- Representation
- Algorithmic solutions
- Transparency as a solution to the fairness problem

## Accountability

- Physical safety – Data safety
- Economic safety

## Transparency

- Explainable AI
- Understanding the inner workings of a superintelligence

# AI ethics, fairness, and bias



Let's define the following random variables:

- $A$  = protected attribute: An observable fact that should not be predictive of outcomes, e.g., gender, race, age, disability.
- $X$  = observable data that we can use for our decision
- $Y$  = the unknown correct label for this person (e.g.,  $Y = 1$  might mean “this person should receive a loan” or “should be admitted to USC”)
- $\hat{Y}$  = a function of  $X$ , designed using probabilistic or neural methods to approximate  $Y$  as closely as possible

# AI ethics, fairness, and bias



## Demographic Parity:

The probability of a positive outcome is the same, regardless of protected attribute.

$$P(\hat{Y} = 1 | A = a) = P(\hat{Y} = 1 | A = a') \quad \forall a, a'$$

## Predictive Parity:

Precision is the same, regardless of protected attribute.

$$P(Y = 1 | \hat{Y} = 1, A = a) = P(Y = 1 | \hat{Y} = 1, A = a') \quad \forall a, a'$$

## Error Rate Balance:

Recall is the same, regardless of protected attribute.

$$P(\hat{Y} = 1 | Y = 1, A = a) = P(\hat{Y} = 1 | Y = 1, A = a') \quad \forall \hat{y}, a, a'$$

# AI ethics, fairness, and bias

Predictive parity

## You can't have all three

From Bayes' theorem  
 $P(M|D) = P(D|M)P(M)/P(D)$

Balanced error

$$P(\hat{Y} = 1|Y = 1, A = a) = \frac{P(Y = 1|\hat{Y} = 1, A = a)P(\hat{Y} = 1|A = a)}{P(Y = 1|A = a)}$$

Demographic parity

Ground Truth

$$P(\hat{Y} = 1|Y = 1, A = a') = \frac{P(Y = 1|\hat{Y} = 1, A = a')P(\hat{Y} = 1|A = a')}{P(Y = 1|A = a')}$$

The balanced error, predictive parity, and demographic parity terms cannot all be independent of A unless Y is also independent of A.

In other words, if the current state of society is unfair (distribution of positive outcomes currently depends on protected attribute), then algorithmic solutions cannot make it fair (at least not in all three ways, all at once).

# Other Useful Definitions of Fairness in AI



## **Individual Fairness:**

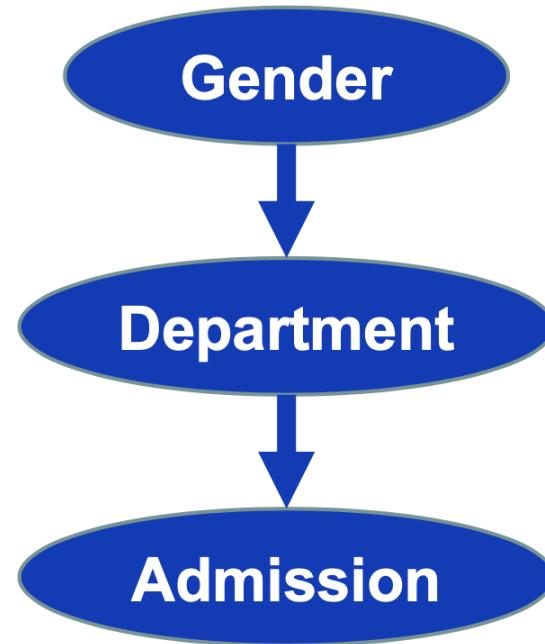
The dissimilarity between two outcomes should be less than the dissimilarity between the people.

## **Counterfactual Fairness:**

If a person's protected attribute were changed (and all their other attributes were possibly changed, according to their dependence on the protected attribute), then the outcome should not change.

# State-of-the-art Solution: Explain Your Assumptions to Your Users

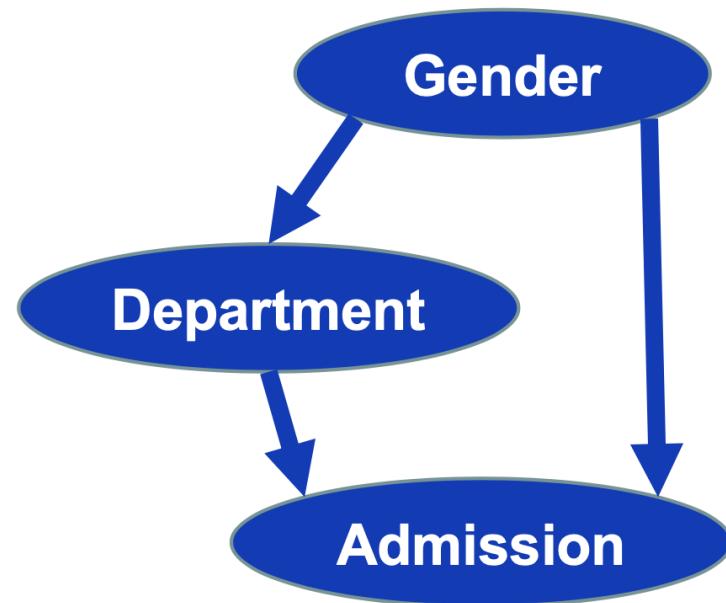
- The Transparency Dogma says that an algorithm's assumptions should be published in a way that users can understand, so that users can participate in a public debate about the fairness of the assumptions.
- A common way to do this is by drawing a Bayesian network.
- Example shown: the Bickel-Hammel-O'Connell data. Admission was conditionally independent of Gender, given Department.



# State-of-the-art Solution: Explain Your Assumptions to Your Users

A possible solution: in order to better approximate demographic parity,

- Make admission explicitly dependent on gender.
- Admit women at a slightly higher percentage, in every department to which they apply, so that...
- ...the total percentage of admitted women equals the total percentage of admitted men.



Public debate: is that more fair, or less fair?